




WORKSHOP

INTRODUCTION TO DEEP LEARNING

Tue Vu, PhD
Research & Data Science Services
SMU OIT

Outline

1. Convolution Neural Network
2. Recurrent Neural Network
3. Long-Short Term Memory



1. Convolution Neural Network


- CNNs are one type of ANN which utilize the neuron, kernel, activation function.
- Inputs **must be** in images (or assumed to be images in 2D format)
- Using Forward & Backpropagation technique with certain property to process it faster
- CNNs best for object detection, image classification, computer vision



1. Convolution Neural Network

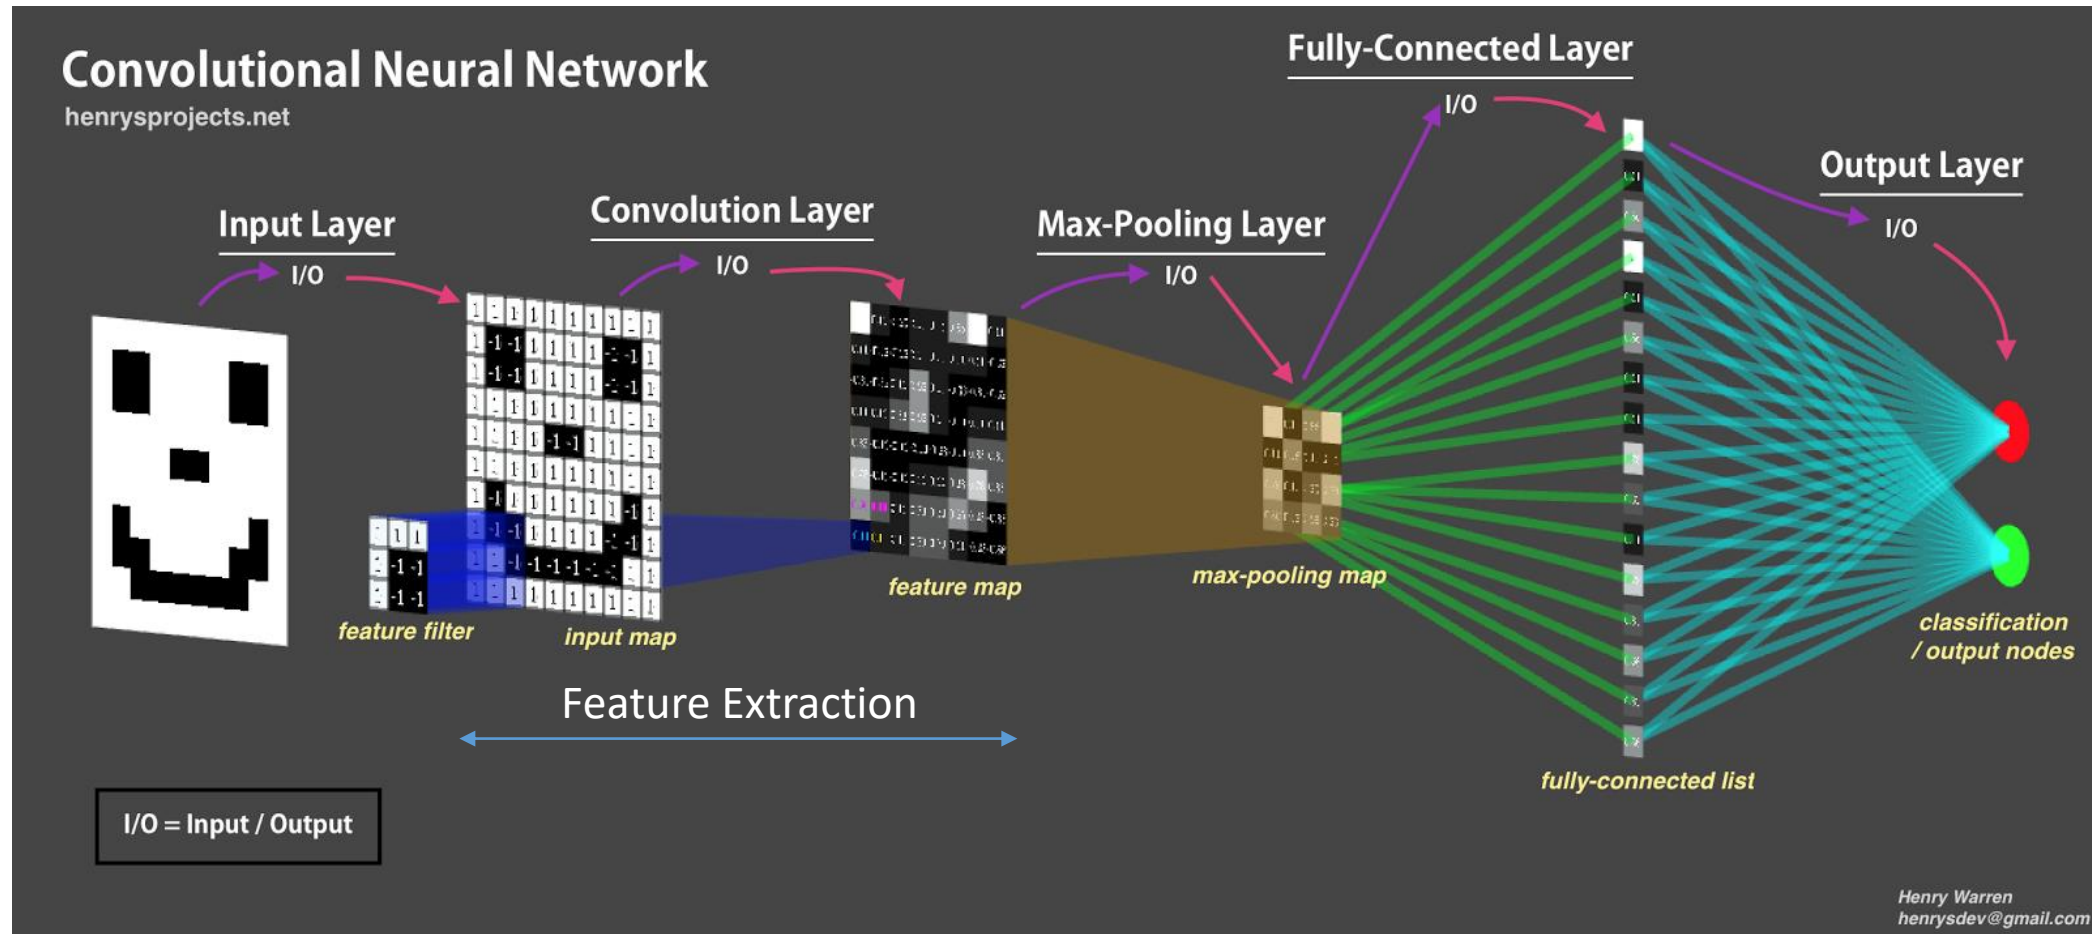
Architecture of MLPs

- Previous examples (MNIST, Fashion MNIST) use fully connected MLP NN to predict the images
- The accuracy/loss are ok but not so great



1. Convolution Neural Network

Architecture of CNNs



1. Convolution Neural Network

Architecture of CNNs

- Convolutional Layers
- Pooling Layers
- Flatten Layer

1. Convolution Neural Network


Parameters of Convolutional Layers (Conv2D):

- Depth
- Filter/kernel
- Stride
- Padding

1. Convolution Neural Network

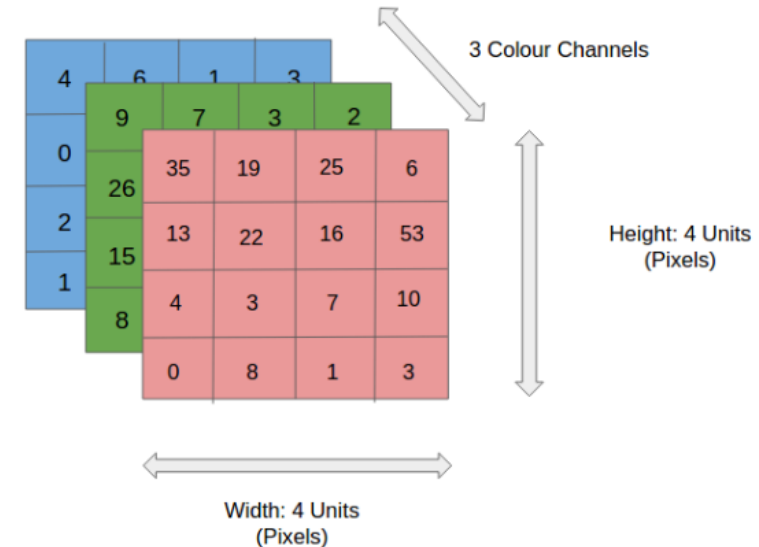
Parameters of Conv2D: depth

Depth = 1



MNIST
Fashion MNIST

**Depth = 3
(RGB)**



Regular
images

1. Convolution Neural Network

Parameters of Conv2D: filter & kernel

1. Convolution Neural Network

Parameters of Conv2D: filter & kernel

- dot product

| Input | | Kernel | | Output | | | | | | | | | | | | | | | | | |
|---|----|--------|---|--------|---|---|---|---|---|---|--|---|---|---|---|---|--|----|----|----|----|
| <table border="1"><tr><td>0</td><td>1</td><td>2</td></tr><tr><td>3</td><td>4</td><td>5</td></tr><tr><td>6</td><td>7</td><td>8</td></tr></table> | 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | * | <table border="1"><tr><td>0</td><td>1</td></tr><tr><td>2</td><td>3</td></tr></table> | 0 | 1 | 2 | 3 | = | <table border="1"><tr><td>19</td><td>25</td></tr><tr><td>37</td><td>43</td></tr></table> | 19 | 25 | 37 | 43 |
| 0 | 1 | 2 | | | | | | | | | | | | | | | | | | | |
| 3 | 4 | 5 | | | | | | | | | | | | | | | | | | | |
| 6 | 7 | 8 | | | | | | | | | | | | | | | | | | | |
| 0 | 1 | | | | | | | | | | | | | | | | | | | | |
| 2 | 3 | | | | | | | | | | | | | | | | | | | | |
| 19 | 25 | | | | | | | | | | | | | | | | | | | | |
| 37 | 43 | | | | | | | | | | | | | | | | | | | | |

1. Convolution Neural Network

Parameters of Conv2D : filter & kernel

Kernel size (3,3)

| | | |
|---------------|---------------|---------------|
| $\frac{1}{9}$ | $\frac{1}{9}$ | $\frac{1}{9}$ |
| $\frac{1}{9}$ | $\frac{1}{9}$ | $\frac{1}{9}$ |
| $\frac{1}{9}$ | $\frac{1}{9}$ | $\frac{1}{9}$ |

Blur filter



1. Convolution Neural Network

Parameters of Conv2D : filter & kernel

Kernel size (3,3)

| | | |
|----|----|----|
| 0 | -1 | 0 |
| -1 | 5 | -1 |
| 0 | -1 | 0 |

Sharp filter



1. Convolution Neural Network

Parameters of Conv2D : filter & kernel

Kernel size (3,3)


| | | |
|----|----|----|
| -1 | -2 | -1 |
| 0 | 0 | 0 |
| 1 | 2 | 1 |

Horizontal

| | | |
|----|---|---|
| -1 | 0 | 1 |
| -2 | 0 | 2 |
| -1 | 0 | 1 |

Vertical

Edge detection



1. Convolution Neural Network

Parameters of Conv2D : filter & kernel

Convolved Feature with filter

- CNN uses the Convolved Feature to reduce the image size by dot product with given kernel (filter)
- The image reduction without losing features and easier to process for good prediction
- In CNNs, filters are not defined. The value of each filter is learned during the training process.

| | | | | |
|-----------------|-----------------|-----------------|---|---|
| 1 _{x1} | 1 _{x0} | 1 _{x1} | 0 | 0 |
| 0 _{x0} | 1 _{x1} | 1 _{x0} | 1 | 0 |
| 0 _{x1} | 0 _{x0} | 1 _{x1} | 1 | 1 |
| 0 | 0 | 1 | 1 | 0 |
| 0 | 1 | 1 | 0 | 0 |

Image

| | | |
|---|--|--|
| 4 | | |
| | | |
| | | |

Convolved
Feature

1. Convolution Neural Network

Parameters of Conv2D : filter & kernel

| | | | | | | |
|-----|-----|-----|-----|-----|-----|-----|
| 0 | 0 | 0 | 0 | 0 | 0 | ... |
| 0 | 156 | 155 | 156 | 158 | 158 | ... |
| 0 | 153 | 154 | 157 | 159 | 159 | ... |
| 0 | 149 | 151 | 155 | 158 | 159 | ... |
| 0 | 146 | 146 | 149 | 153 | 158 | ... |
| 0 | 145 | 143 | 143 | 148 | 158 | ... |
| ... | ... | ... | ... | ... | ... | ... |

| | | | | | | |
|-----|-----|-----|-----|-----|-----|-----|
| 0 | 0 | 0 | 0 | 0 | 0 | ... |
| 0 | 167 | 166 | 167 | 169 | 169 | ... |
| 0 | 164 | 165 | 168 | 170 | 170 | ... |
| 0 | 160 | 162 | 166 | 169 | 170 | ... |
| 0 | 156 | 156 | 159 | 163 | 168 | ... |
| 0 | 155 | 153 | 153 | 158 | 168 | ... |
| ... | ... | ... | ... | ... | ... | ... |

| | | | | | | |
|-----|-----|-----|-----|-----|-----|-----|
| 0 | 0 | 0 | 0 | 0 | 0 | ... |
| 0 | 163 | 162 | 163 | 165 | 165 | ... |
| 0 | 160 | 161 | 164 | 166 | 166 | ... |
| 0 | 156 | 158 | 162 | 165 | 166 | ... |
| 0 | 155 | 155 | 158 | 162 | 167 | ... |
| 0 | 154 | 152 | 152 | 157 | 167 | ... |
| ... | ... | ... | ... | ... | ... | ... |

Input Channel #1 (Red)

Input Channel #2 (Green)

Input Channel #3 (Blue)

Kernel size (3,3)

| | | |
|----|----|----|
| -1 | -1 | 1 |
| 0 | 1 | -1 |
| 0 | 1 | 1 |

Kernel Channel #1

| | | |
|---|----|----|
| 1 | 0 | 0 |
| 1 | -1 | -1 |
| 1 | 0 | -1 |

Kernel Channel #2

| | | |
|---|----|---|
| 0 | 1 | 1 |
| 0 | 1 | 0 |
| 1 | -1 | 1 |

Kernel Channel #3

3 kernels = 3 filters

308

+

-498

+

164

+ 1 = -25

Bias = 1


Output

| | | | | |
|-----|-----|-----|-----|-----|
| -25 | | | | ... |
| | | | | ... |
| | | | | ... |
| | | | | ... |
| ... | ... | ... | ... | ... |

1. Convolution Neural Network

Parameters of Conv2D : stride

- Stride tuned for the compression of images and video data




1. Convolution Neural Network

Parameters of Conv2D : stride


- Stride tuned for the compression of images and video data

Stride = (1,1)



Image

Stride = (2,2)




Image


1. Convolution Neural Network

Parameters of Conv2D : stride


- Stride tuned for the compression of images and video data



Image



Image



Image

1. Convolution Neural Network

Parameters of Conv2D : padding

- The pixels located on the corners and the edges are used much less than those in the middle => the information on borders and edges are not preserved

| | | | | |
|-----------------|-----------------|-----------------|---|---|
| 1 _{x1} | 1 _{x0} | 1 _{x1} | 0 | 0 |
| 0 _{x0} | 1 _{x1} | 1 _{x0} | 1 | 0 |
| 0 _{x1} | 0 _{x0} | 1 _{x1} | 1 | 1 |
| 0 | 0 | 1 | 1 | 0 |
| 0 | 1 | 1 | 0 | 0 |

Image

| | | |
|---|--|--|
| 4 | | |
| | | |
| | | |

Convolved
Feature

1. Convolution Neural Network

Parameters of Conv2D : padding

- The pixels located on the corners and the edges are used much less than those in the middle => the information on borders and edges are not preserved

| | | | | |
|-----------------|-----------------|-----------------|---|---|
| 1 _{x1} | 1 _{x0} | 1 _{x1} | 0 | 0 |
| 0 _{x0} | 1 _{x1} | 1 _{x0} | 1 | 0 |
| 0 _{x1} | 0 _{x0} | 1 _{x1} | 1 | 1 |
| 0 | 0 | 1 | 1 | 0 |
| 0 | 1 | 1 | 0 | 0 |

Image

| | | |
|---|--|--|
| 4 | | |
| | | |
| | | |

Convolved
Feature

Solution?

1. Convolution Neural Network

Parameters of Conv2D : padding

- The pixels located on the corners and the edges are used much less than those in the middle => the information on borders and edges are not preserved

| | | | | |
|----|----|---|----|----|
| 3 | 5 | 9 | 1 | 10 |
| 13 | 2 | 4 | 6 | 11 |
| 16 | 24 | 9 | 13 | 1 |
| 7 | 1 | 6 | 8 | 3 |
| 8 | 4 | 9 | 1 | 9 |

padding →

| | | | | | | | | |
|---|---|----|----|---|----|----|---|---|
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 3 | 5 | 9 | 1 | 10 | 0 | 0 |
| 0 | 0 | 13 | 2 | 4 | 6 | 11 | 0 | 0 |
| 0 | 0 | 16 | 24 | 9 | 13 | 1 | 0 | 0 |
| 0 | 0 | 7 | 1 | 6 | 8 | 3 | 0 | 0 |
| 0 | 0 | 8 | 4 | 9 | 1 | 9 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |

Solution?

- Add rows and columns of 0 to the input images
- The image on left was added with padding parameter $P=2$

(W, H)

(W + 2P, H + 2P)

1. Convolution Neural Network

Parameters of Conv2D : padding

- The pixels located on the corners and the edges are used much less than those in the middle => the information on borders and edges are not preserved

| | | | | |
|----|----|---|----|----|
| 3 | 5 | 9 | 1 | 10 |
| 13 | 2 | 4 | 6 | 11 |
| 16 | 24 | 9 | 13 | 1 |
| 7 | 1 | 6 | 8 | 3 |
| 8 | 4 | 9 | 1 | 9 |

padding →

| | | | | | | | | |
|---|---|----|----|---|----|----|---|---|
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 3 | 5 | 9 | 1 | 10 | 0 | 0 |
| 0 | 0 | 13 | 2 | 4 | 6 | 11 | 0 | 0 |
| 0 | 0 | 16 | 24 | 9 | 13 | 1 | 0 | 0 |
| 0 | 0 | 7 | 1 | 6 | 8 | 3 | 0 | 0 |
| 0 | 0 | 8 | 4 | 9 | 1 | 9 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |

In Keras

- padding = "valid": no padding
- padding = "same": padding with 0 evenly left/right, up/down
- padding = "same" with strides = (1,1): output has same size as input

(W, H)

(W + 2P, H + 2P)

1. Convolution Neural Network

Parameters of Conv2D :

- Depth (L): 3 or 1
- Number of Filter: F
- kernel: (K, K): (3,3) or (5,5)
- Stride (S, S): (1,1) or (2,2)
- Padding (P)_w

Formulation to compute the output size of a convolutional layer from an image with size (W, H)?

| | | | | | | | |
|---|-----|-----|-----|-----|-----|-----|-----|
| H | 0 | 0 | 0 | 0 | 0 | 0 | ... |
| | 0 | 156 | 155 | 156 | 158 | 158 | ... |
| | 0 | 153 | 154 | 157 | 159 | 159 | ... |
| | 0 | 149 | 151 | 155 | 158 | 159 | ... |
| | 0 | 146 | 146 | 149 | 153 | 158 | ... |
| | 0 | 145 | 143 | 143 | 148 | 158 | ... |
| | ... | ... | ... | ... | ... | ... | ... |


Input Channel #1 ($S=1$)

1. Convolution Neural Network

Parameters of Conv2D :

- Depth (L): 3 or 1
- Number of Filter: F
- kernel: (K, K): (3,3) or (5,5)
- Stride (S, S): (1,1) or (2,2)
- Padding (P)

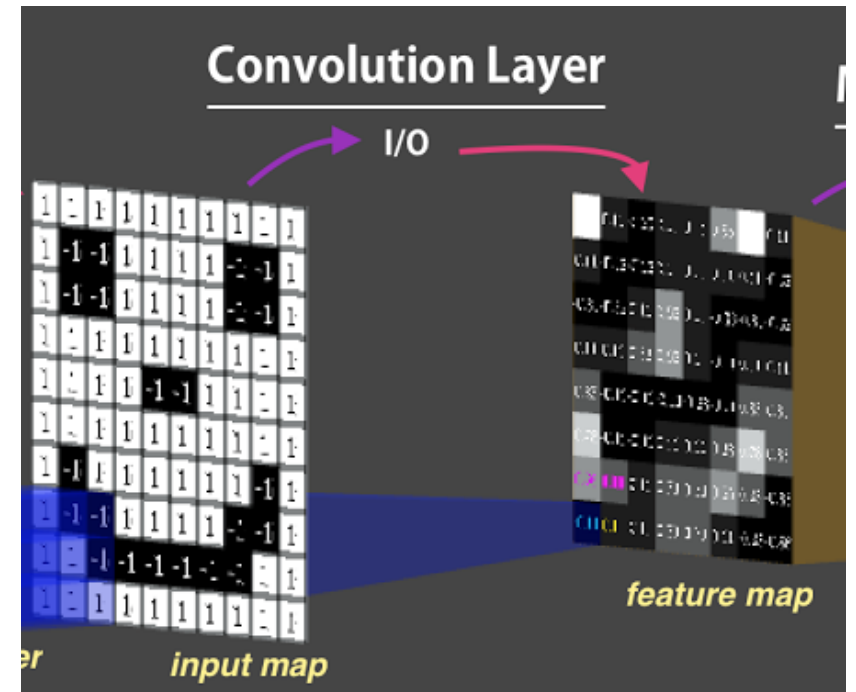
Formulation to compute the output size of a convolutional layer from an image with size (W, H)?



1. Convolution Neural Network

How to add Conv2D in keras?


```
model = Sequential()  
model.add(Conv2D(F, (K, K), strides=(S, S), activation='relu', padding="same", input_shape=(32, 32, L)))
```



1. Convolution Neural Network

Convolutional Layer (CNN or ConvNet):

- The CNN will reduce the original RGB images to its Convolutional Layer
- Multiple layers can be applied



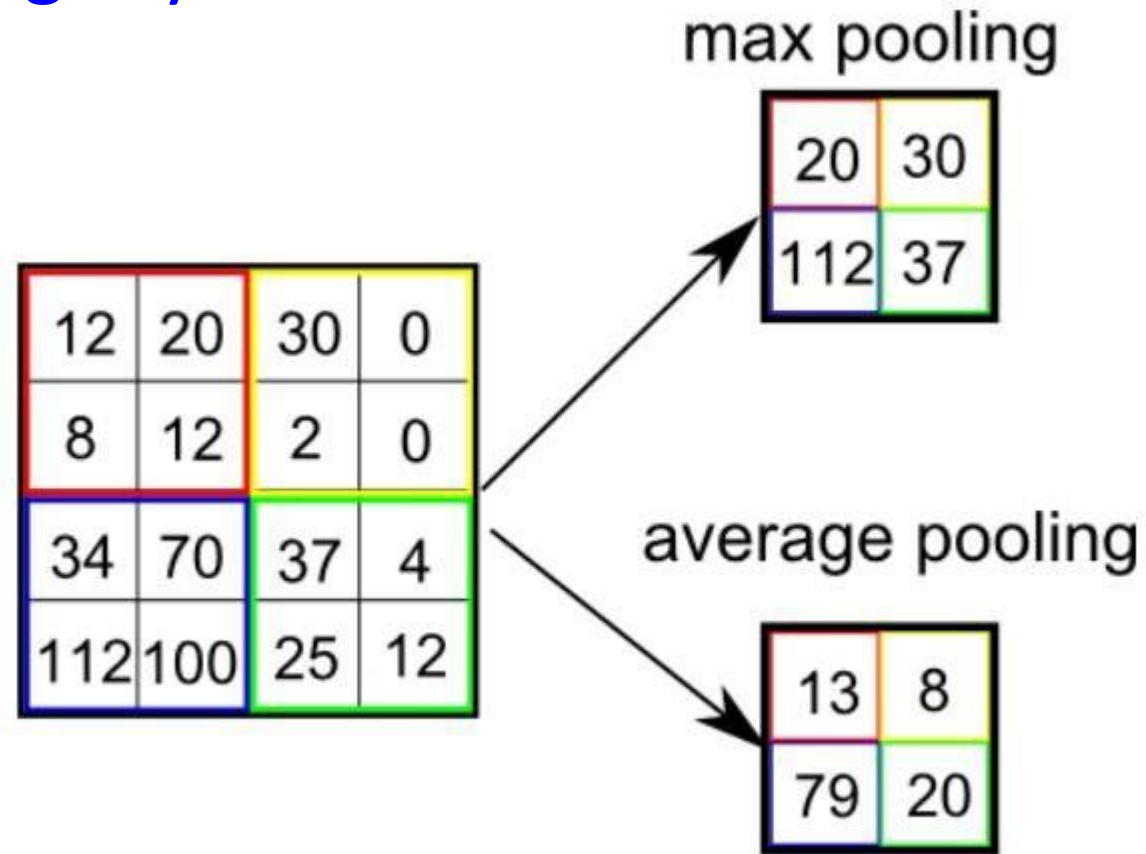
1. Convolution Neural Network

Pooling Layer

- Pooling Layer should follow Convolutional Layer
- Similar to the Convolutional Layer, the Pooling layer is responsible for reducing the spatial size of the Convolved Feature.
- This is to decrease the computational power required to process the data through dimensionality reduction
- Two types of Pooling: Max Pooling & Average Pooling.

1. Convolution Neural Network

Pooling Layer




In which Max Pooling performs a lot better than Average Pooling.

1. Convolution Neural Network

Pooling Layer


- The image after Max Pooling layer would look like:



1. Convolution Neural Network


Flatten Layer

- Once the images have passed through Convolution Layer and Pooling Layer, its size has been reduced greatly and ready for MLP training (or to another Convolution steps).
- The image is then flatten to a column vector and passed through feed-forward NN and BackPropagation applied to every iteration.
- Softmax activation function is applied to classified the multi-output/multi-labels



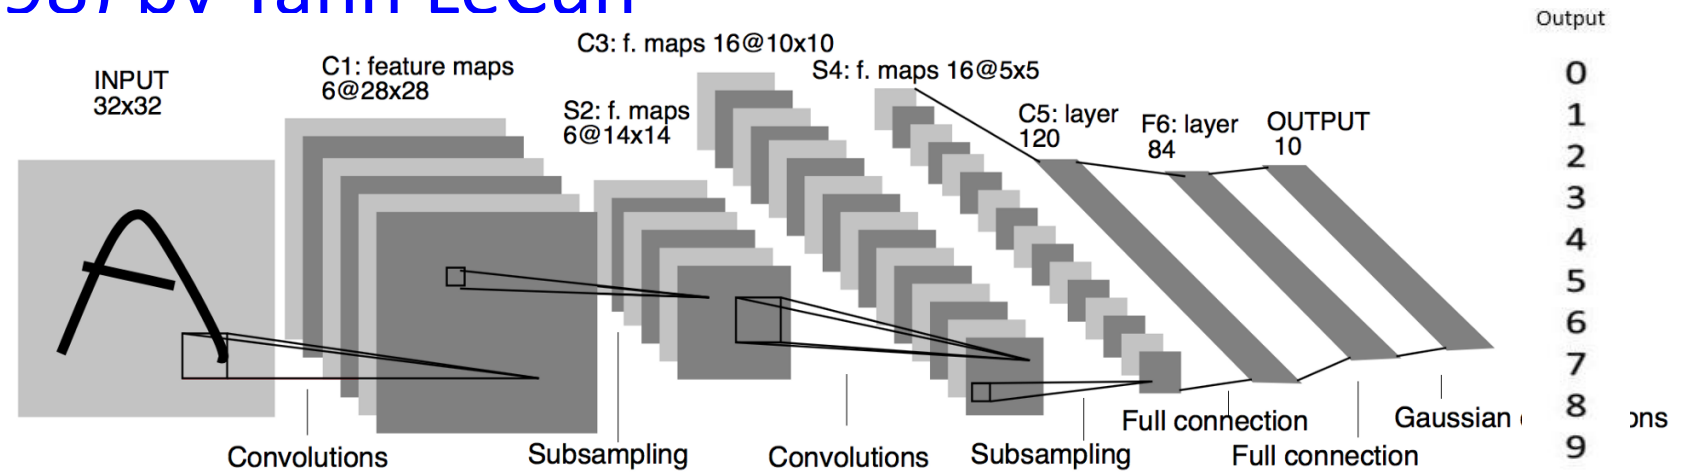
1. Convolution Neural Network

A Sample of CNN:



1. Convolution Neural Network

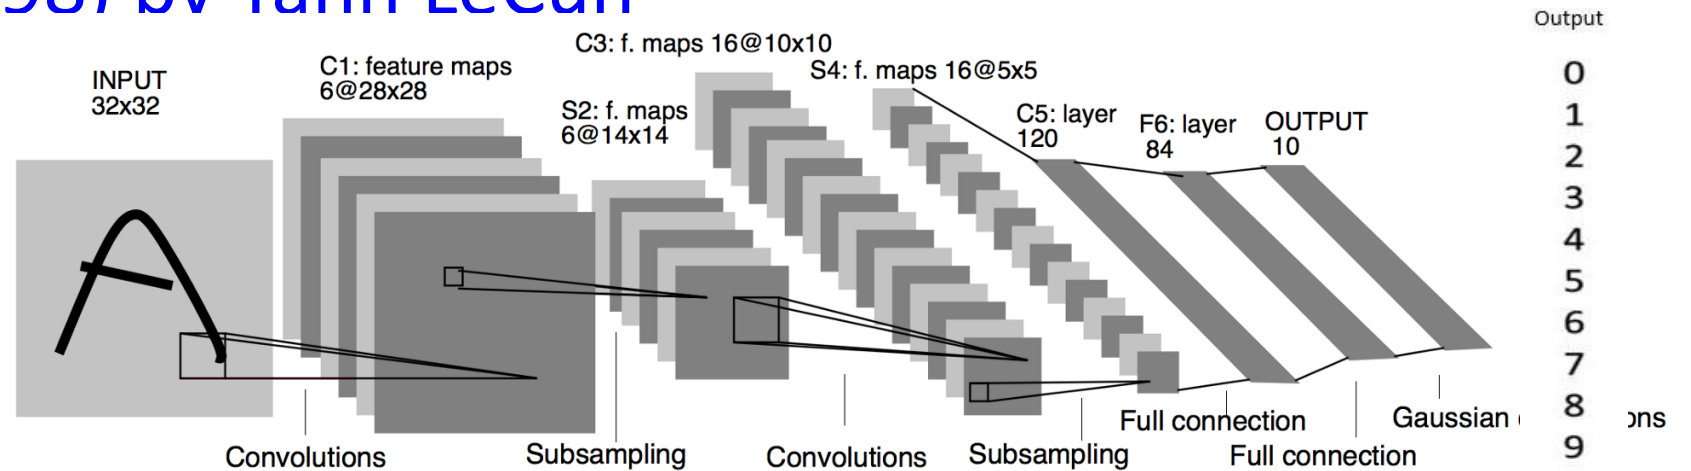
Lenet-5 (1998) by Yann LeCun



- LeNet-5 is designed for handwritten and machine-printed character recognition
- Input of 32x32x1
- Total parameters: 60k
- Activation function: tanh

1. Convolution Neural Network

Lenet-5 (1998) by Yann LeCun

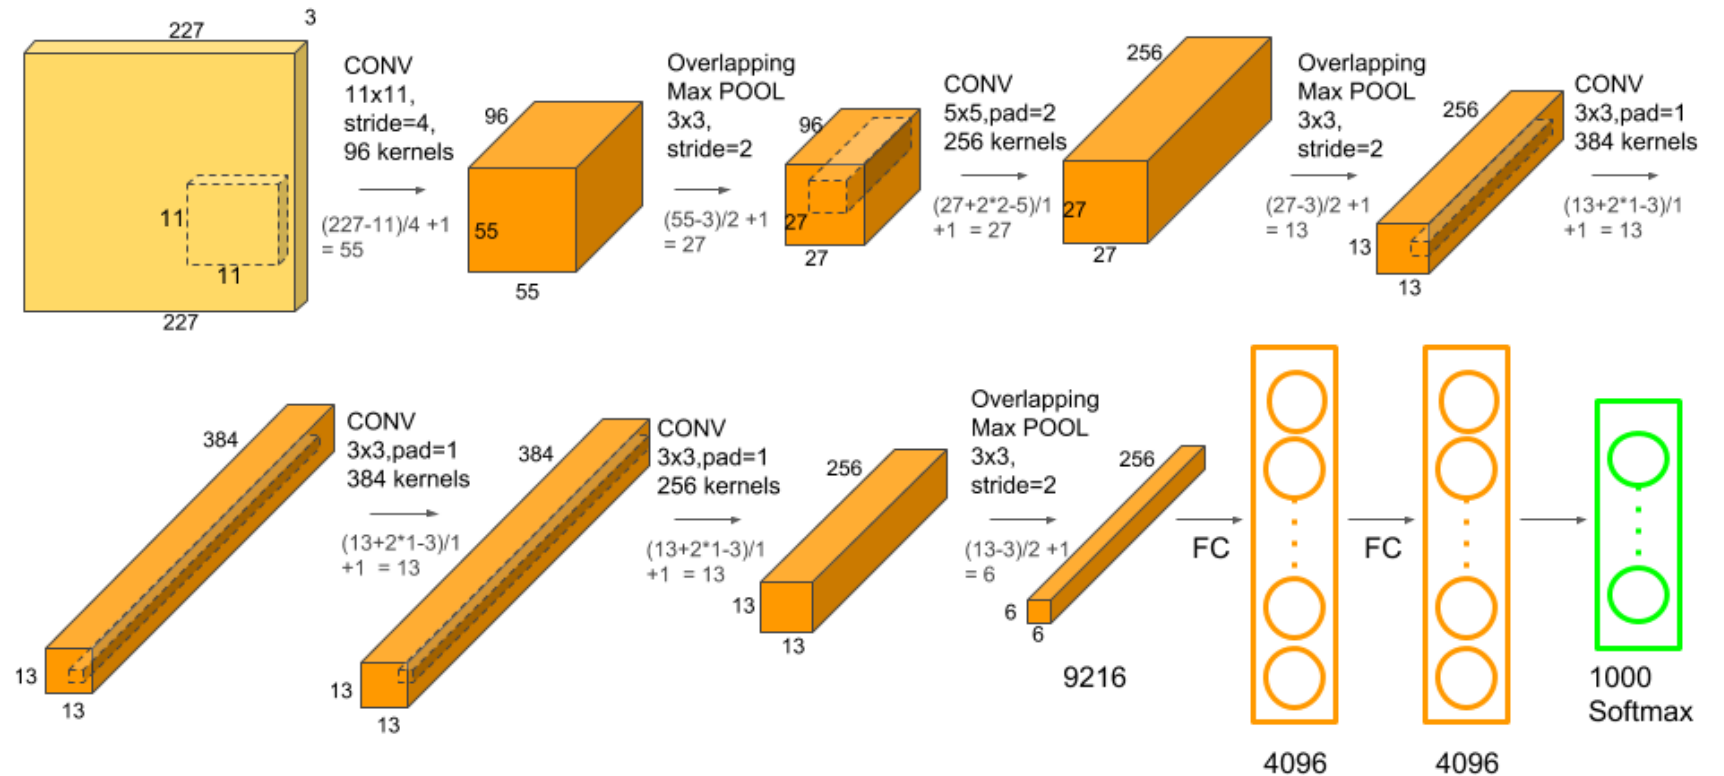


```
model = Sequential()
model.add(Conv2D(6, (5, 5), strides=(1, 1), activation='tanh', padding="valid", input_shape=(32, 32, 1)))
model.add(AveragePooling2D(pool_size=(2, 2), strides=(2, 2)))
model.add(Conv2D(16, (5, 5), strides=(1, 1), activation='tanh', padding="valid"))
model.add(AveragePooling2D(pool_size=(2, 2), strides=(2, 2)))
model.add(Conv2D(120, (5, 5), strides=(1, 1), activation='tanh', padding="valid"))
model.add(Flatten())
model.add(Dense(84, activation='tanh'))
model.add(Dense(10, activation='softmax'))
```

1. Convolution Neural Network

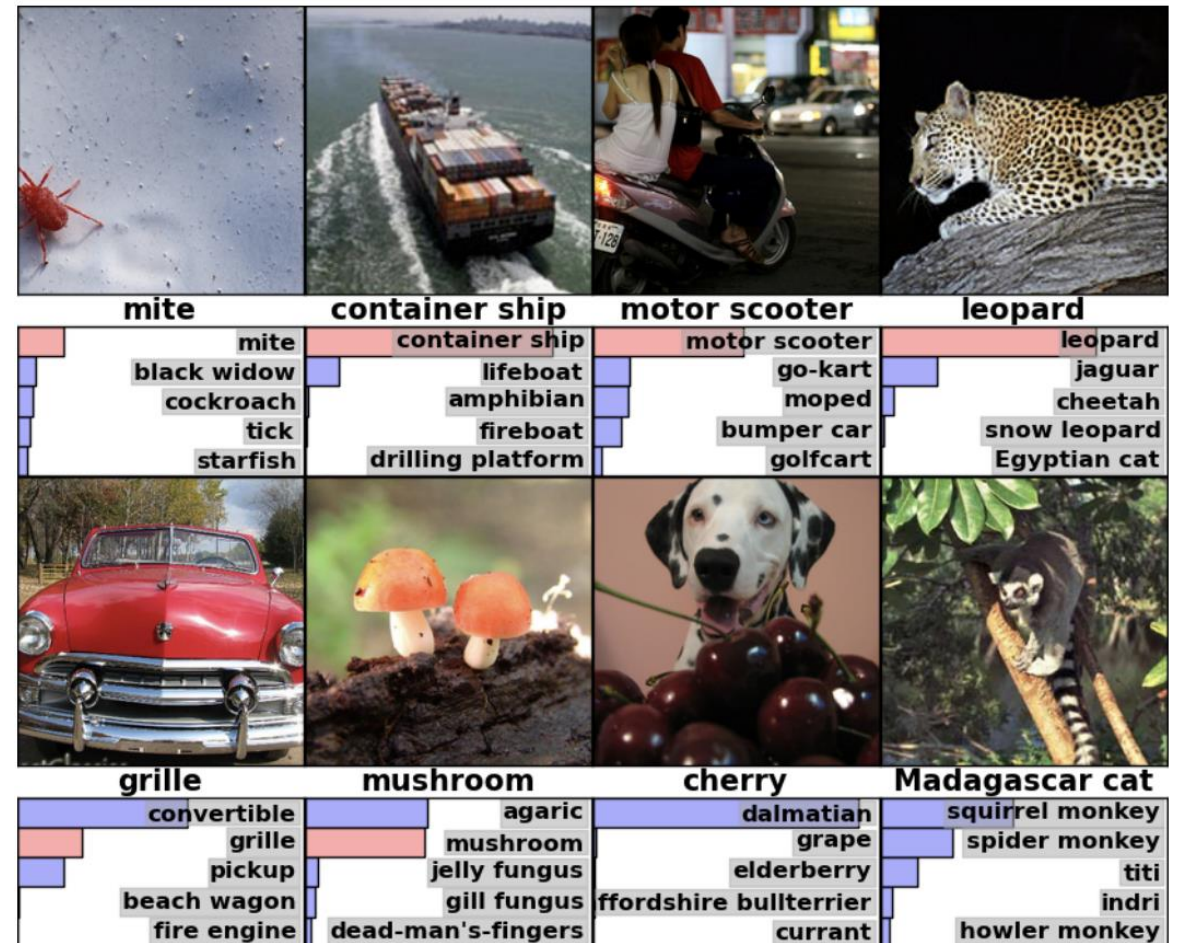
AlexNet (2012) by Hinton Alex Krizhevsky

- AlexNet won the 2012 ImageNet challenge
- Input of 227x227x3
- Total parameters: 60M
- Activation: ReLU



1. Convolution Neural Network


AlexNet (2012) by Hinton Alex Krizhevsky



1. Convolution Neural Network

VGG16 (2014) - Visual Geometry Group


- VGG16 runner up of 2014 ImageNet challenge
- 16 layers: 13 ConvNet, 3 Fully Connected
- Total Parameters: 130M



1. Convolution Neural Network

GoogleNet (2014)

- GoogleNet won the 2014 ImageNet challenge
- Introduced Inception Network
- 22 layers deep with 27 pooling layers and 9 inception models




2. Recurrent Neural Network

2. Recurrent Neural Network

Introduction

- RNNs are type of Deep Learning models with built-in feedback mechanism.
- The output of a particular layer can be **re-fed** as the input in order to predict the output.
- This is different from traditional ML where output/predictand cannot be used as input




Recurrent Neural Networks have loops.

2. Recurrent Neural Network

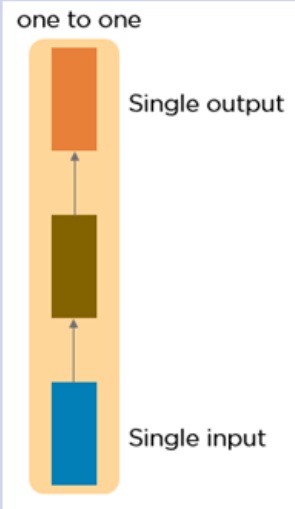
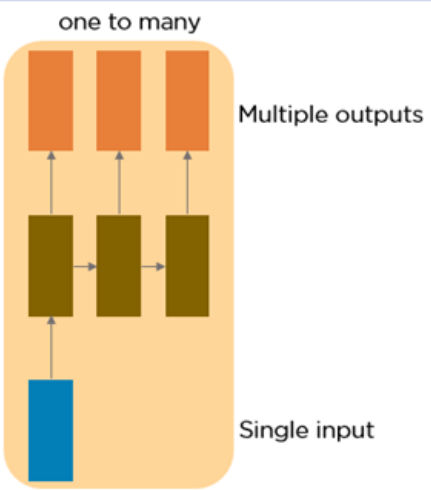
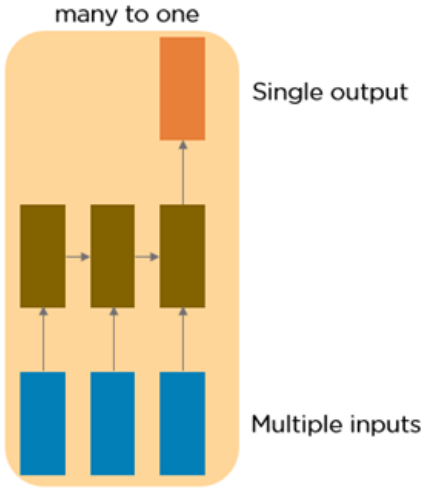
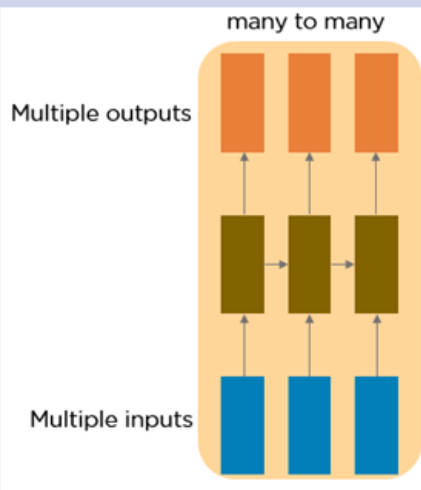


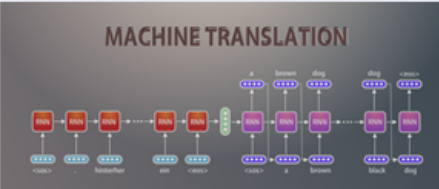
Introduction

- Unroll the RNN loop



2. Recurrent Neural Network

Type of RNNs

| One to One | One to Many | Many to One | Many to Many |
|---|---|---|---|
|  <p>one to one</p> <p>Single output</p> <p>Single input</p> |  <p>one to many</p> <p>Multiple outputs</p> <p>Single input</p> |  <p>many to one</p> <p>Single output</p> <p>Multiple inputs</p> |  <p>many to many</p> <p>Multiple outputs</p> <p>Multiple inputs</p> |
| <p>So called Vanilla NN. Similar to Backpropagation for general ML problem</p> | <p>Image captioning</p>  <p>• large brown dog running away from the sprinkler in the grass . • a brown dog chases the water from a sprinkler on a lawn . • a brown dog running on a lawn near a garden hose • a brown dog plays with the hose . • is about to interact with a fence</p> | <p>Sentiment analysis</p>  | <p>Machine translation</p>  |

2. Recurrent Neural Network

Applications

It is specifically designed for Sequential problem **Weather forecast, Stock forecast, Image captioning, Natural Language Processing, Speech/Voice Recognition**

2. Recurrent Neural Network

Some Disadvantages of RNN:

- Computationally Expensive and large memory requested
- RNN is sensitive to changes in parameters and having gradient problem such as **Exploding Gradient** or **Vanishing Gradient**
- In order to resolve the gradient problem of RNN, a method **Long-Short Term Memory (LSTM)** is proposed.

In this limited workshop, we only cover LSTM for timeseries forecast problem (stock forecast and weather forecast)

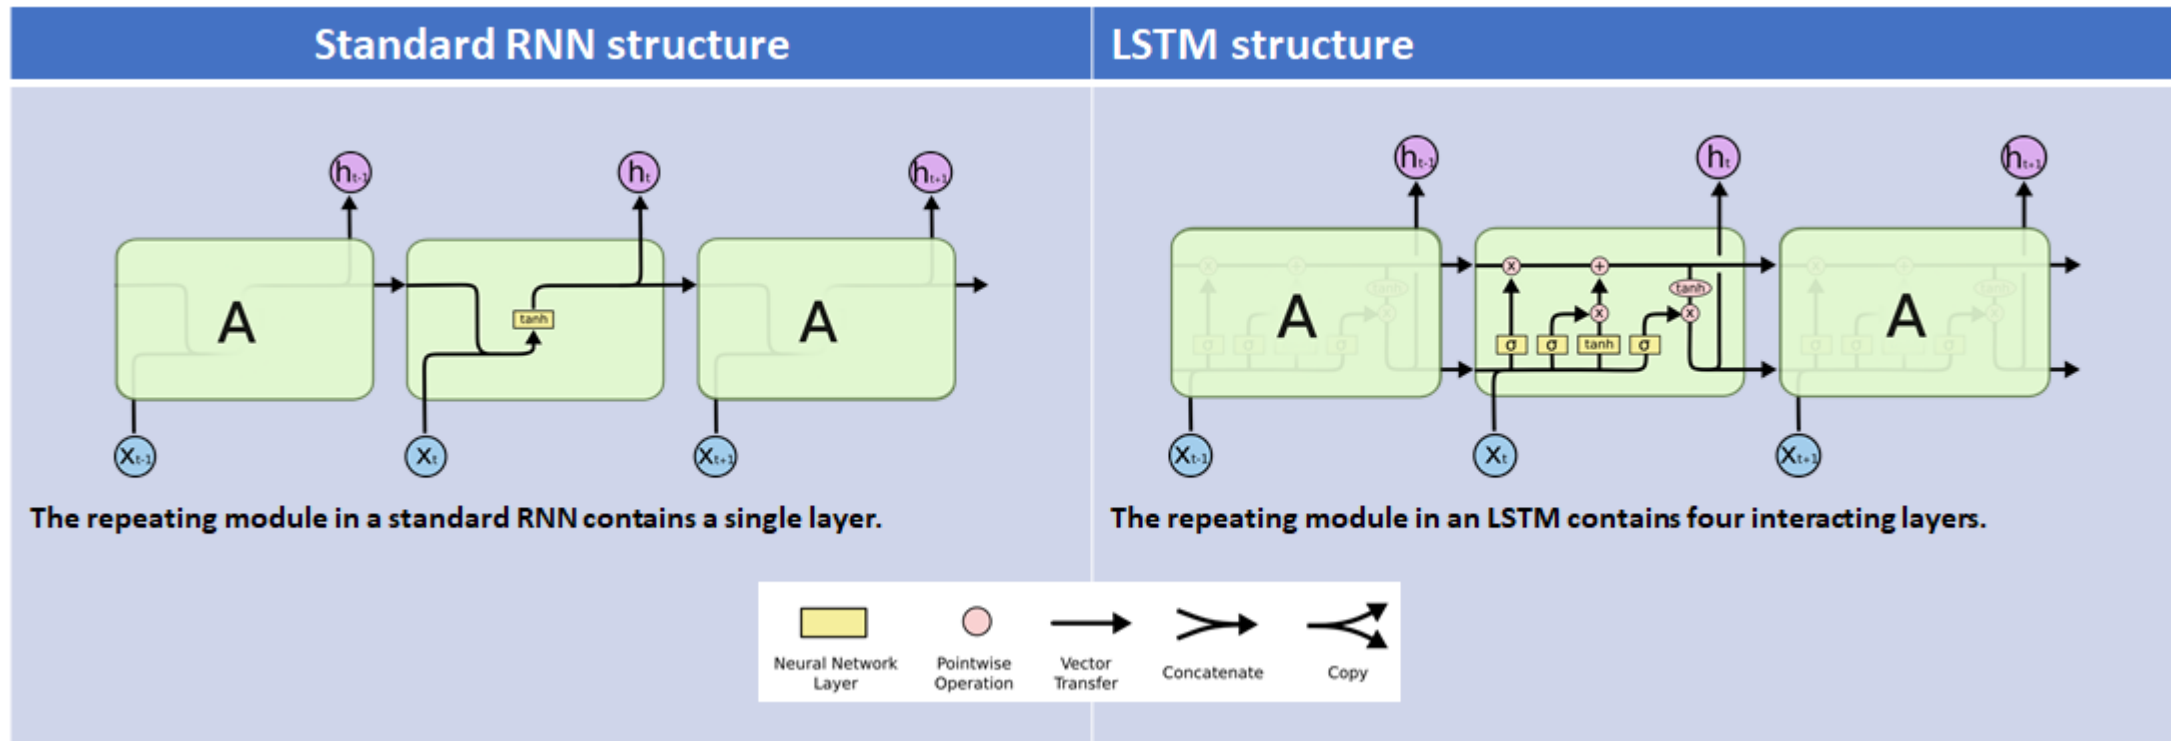
2. Recurrent Neural Network

Long-Short Term Memory model - LSTM

- LSTMs are a special kind of RNN — capable of learning long-term dependencies by remembering information for long periods is the default behavior.
- They were introduced by Hochreiter & Schmidhuber (1997) and were refined and popularized by many people
- LSTMs are explicitly designed to avoid the long-term dependency problem.

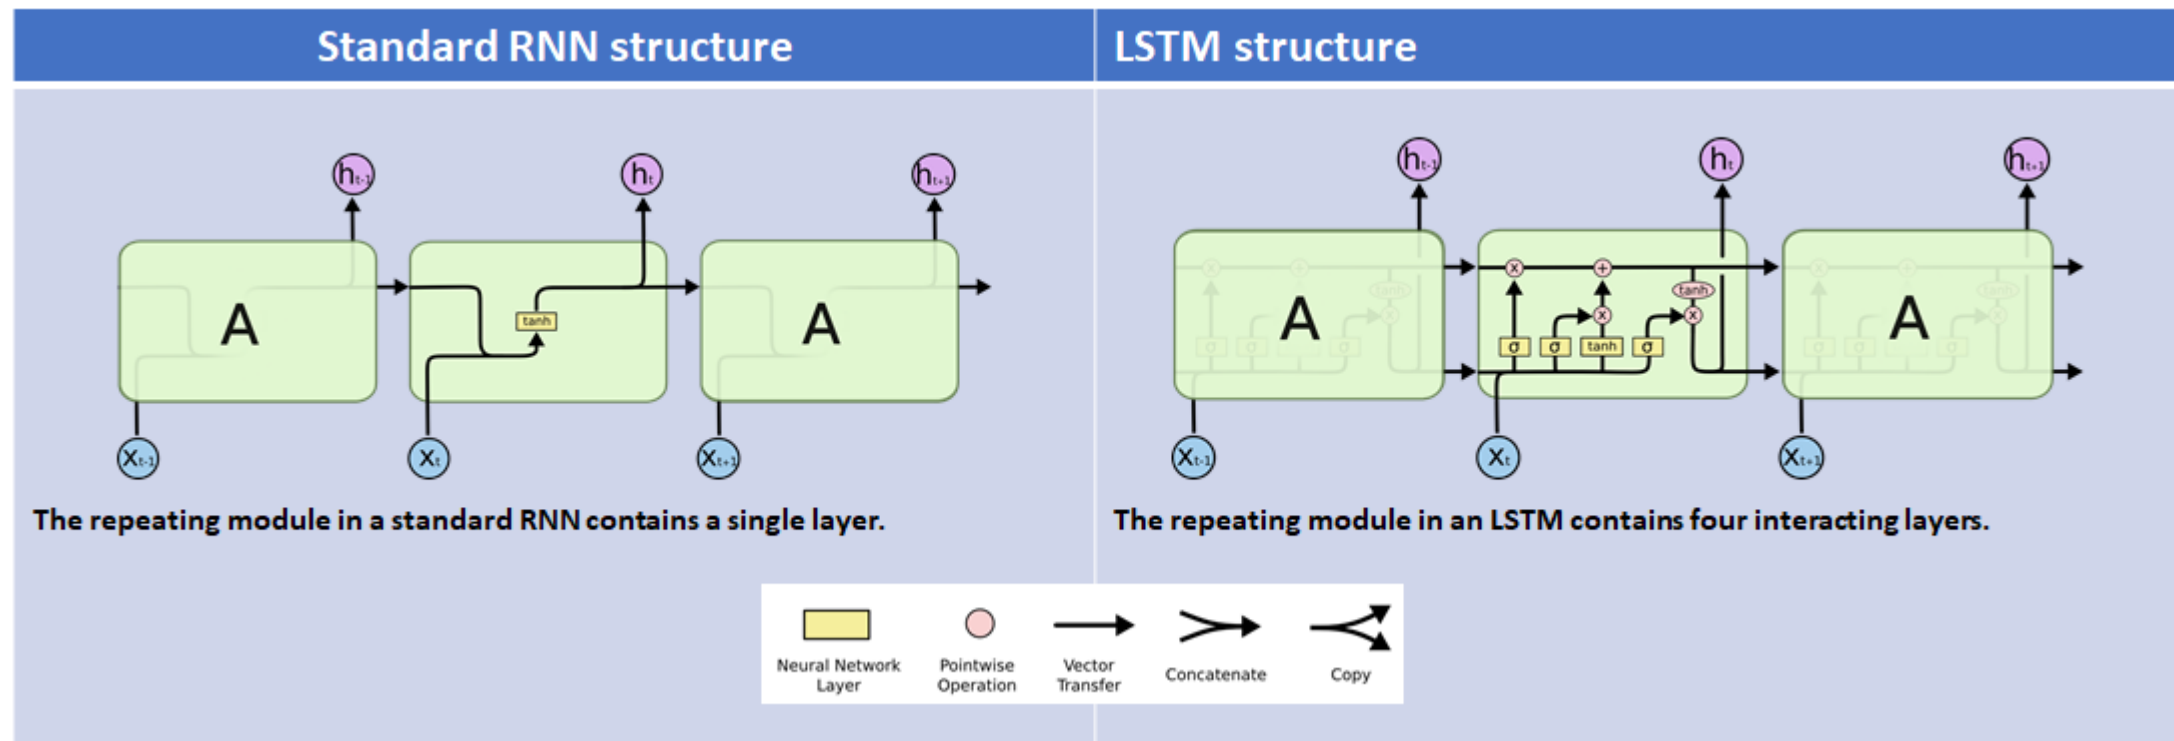
2. Recurrent Neural Network

Comparison between traditional RNN and LSTM



2. Recurrent Neural Network

Comparison between traditional RNN and LSTM



```
model = Sequential()  
model.add(LSTM(128, return_sequences=True, input_shape= (x_train.shape[1], 1)))  
model.add(LSTM(64, return_sequences=False))
```

2. Recurrent Neural Network

Hands-on section