




WORKSHOP

INTRODUCTION TO DEEP LEARNING

Tue Vu, PhD
Research & Data Science Services
SMU OIT


Outline

1. Convolution Neural Network
2. Recurrent Neural Network
3. Long-Short Term Memory



1. Convolution Neural Network

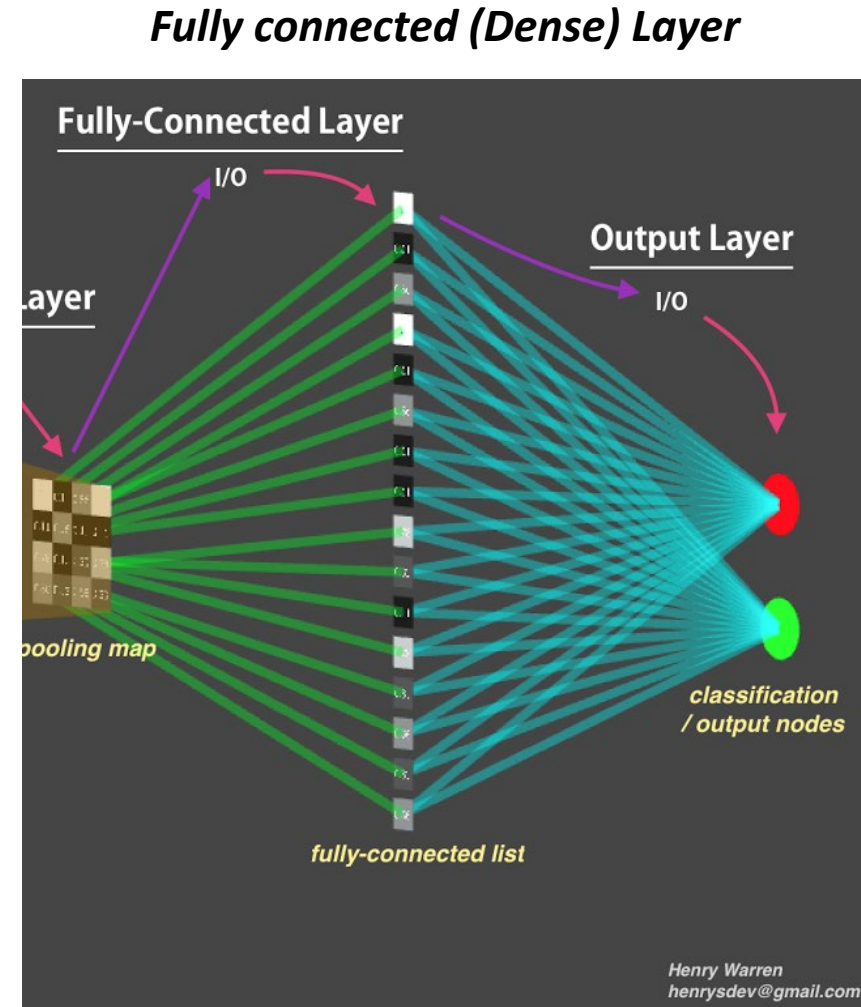
- CNNs are one type of ANN which utilize the neuron, kernel, activation function.
- Inputs **must be** in images (or assumed to be images in 2D format)
- Using Forward & Backpropagation technique with certain property to process it faster
- CNNs best for object detection, image classification, computer vision



1. Convolution Neural Network

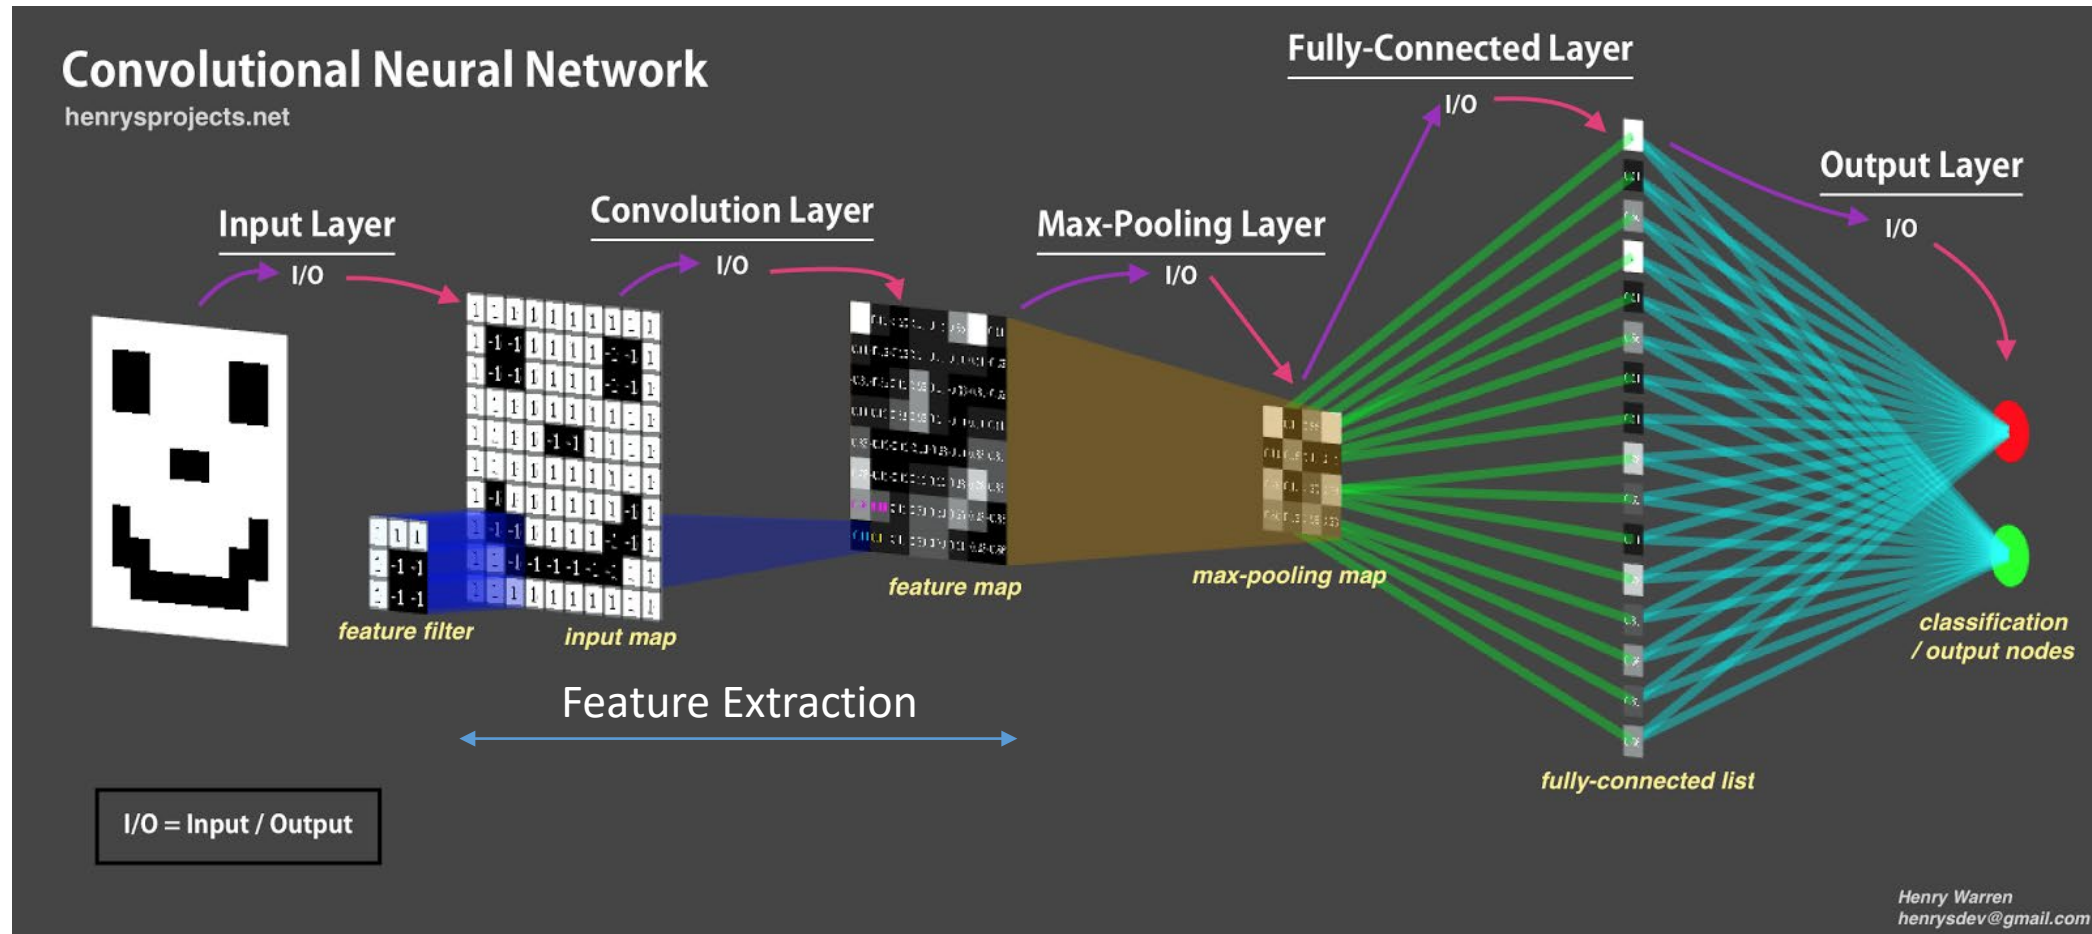
Architecture of MLPs

- Previous examples (MNIST, Fashion MNIST) use fully connected MLP NN to predict the images
- The accuracy/loss are ok but not so great



1. Convolution Neural Network

Architecture of CNNs



1. Convolution Neural Network

Architecture of CNNs

- Convolutional Layers
- Pooling Layers
- Flatten Layer

1. Convolution Neural Network


Hyper-parameters of Convolutional Layers (Conv2D):

- Depth
- Filter/kernel
- Stride
- Padding

1. Convolution Neural Network


Hyper-Parameters of Conv2D: depth

Depth = 1



MNIST
Fashion MNIST

**Depth = 3
(RGB)**



Regular
images

1. Convolution Neural Network

Hyper-Parameters of Conv2D: filter & kernel

1. Convolution Neural Network

Hyper-Parameters of Conv2D: filter & kernel

- dot product

Input		Kernel		Output																	
<table border="1"><tr><td>0</td><td>1</td><td>2</td></tr><tr><td>3</td><td>4</td><td>5</td></tr><tr><td>6</td><td>7</td><td>8</td></tr></table>	0	1	2	3	4	5	6	7	8	*	<table border="1"><tr><td>0</td><td>1</td></tr><tr><td>2</td><td>3</td></tr></table>	0	1	2	3	=	<table border="1"><tr><td>19</td><td>25</td></tr><tr><td>37</td><td>43</td></tr></table>	19	25	37	43
0	1	2																			
3	4	5																			
6	7	8																			
0	1																				
2	3																				
19	25																				
37	43																				


1. Convolution Neural Network

Hyper-Parameters of Conv2D : filter & kernel

Kernel size (3,3)

$\frac{1}{9}$	$\frac{1}{9}$	$\frac{1}{9}$
$\frac{1}{9}$	$\frac{1}{9}$	$\frac{1}{9}$
$\frac{1}{9}$	$\frac{1}{9}$	$\frac{1}{9}$

Blur filter



1. Convolution Neural Network

Hyper-Parameters of Conv2D : filter & kernel

Kernel size (3,3)

0	-1	0
-1	5	-1
0	-1	0

Sharp filter



1. Convolution Neural Network

Hyper-Parameters of Conv2D : filter & kernel

Kernel size (3,3)


-1	-2	-1
0	0	0
1	2	1

Horizontal

-1	0	1
-2	0	2
-1	0	1

Vertical

Edge detection



1. Convolution Neural Network

Hyper-Parameters of Conv2D : filter & kernel

Convolved Feature with filter

- CNN uses the Convolved Feature to reduce the image size by dot product with given kernel (filter)
- The image reduction without losing features and easier to process for good prediction
- In CNNs, filters are not defined. The value of each filter is learned during the training process.

1 _{x1}	1 _{x0}	1 _{x1}	0	0
0 _{x0}	1 _{x1}	1 _{x0}	1	0
0 _{x1}	0 _{x0}	1 _{x1}	1	1
0	0	1	1	0
0	1	1	0	0

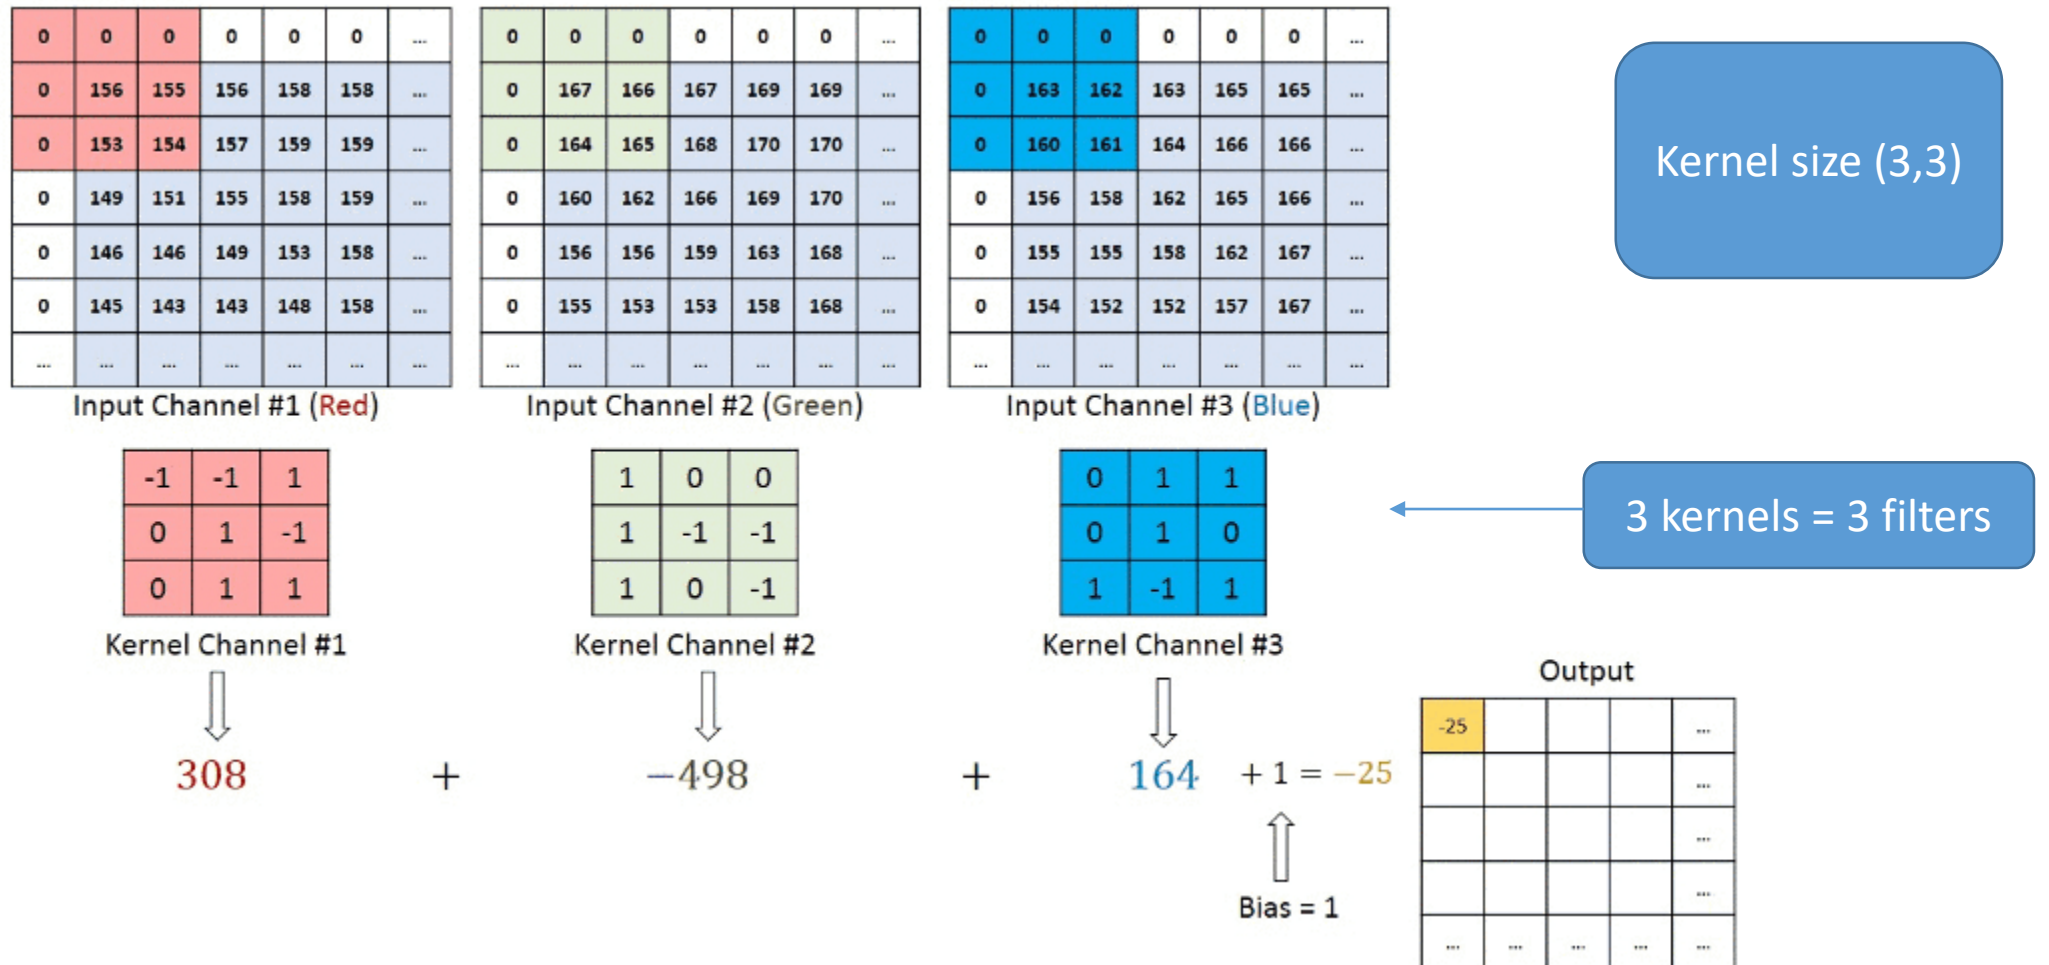
Image

4		

Convolved
Feature

1. Convolution Neural Network


Hyper-Parameters of Conv2D : filter & kernel



1. Convolution Neural Network

Hyper-Parameters of Conv2D : stride

- Stride tuned for the compression of images and video data




1. Convolution Neural Network

Hyper-Parameters of Conv2D : stride


- Stride tuned for the compression of images and video data

Stride = (1,1)



Image

Stride = (2,2)




Image

1. Convolution Neural Network

Hyper-Parameters of Conv2D : stride


- Stride tuned for the compression of images and video data

Stride = (1,1)




Image

Stride = (2,2)



Image

Stride = (2,1) or (1,2)?
Yes but not common



Image

1. Convolution Neural Network

Hyper-Parameters of Conv2D : padding

- The pixels located on the corners and the edges are used much less than those in the middle => the information on borders and edges are not preserved

1 _{x1}	1 _{x0}	1 _{x1}	0	0
0 _{x0}	1 _{x1}	1 _{x0}	1	0
0 _{x1}	0 _{x0}	1 _{x1}	1	1
0	0	1	1	0
0	1	1	0	0

Image

4		

Convolved
Feature

1. Convolution Neural Network

Hyper-Parameters of Conv2D : padding

- The pixels located on the corners and the edges are used much less than those in the middle => the information on borders and edges are not preserved

1 _{x1}	1 _{x0}	1 _{x1}	0	0
0 _{x0}	1 _{x1}	1 _{x0}	1	0
0 _{x1}	0 _{x0}	1 _{x1}	1	1
0	0	1	1	0
0	1	1	0	0

Image

4		

Convolved
Feature

Solution?

1. Convolution Neural Network

Hyper-Parameters of Conv2D : padding

- The pixels located on the corners and the edges are used much less than those in the middle => the information on borders and edges are not preserved



Solution?

- Add rows and columns of 0 to the input images
- The image on left was added with padding parameter $P=2$

(W, H)

(W + 2P, H + 2P)

1. Convolution Neural Network

Hyper-Parameters of Conv2D : padding

- The pixels located on the corners and the edges are used much less than those in the middle => the information on borders and edges are not preserved

3	5	9	1	10
13	2	4	6	11
16	24	9	13	1
7	1	6	8	3
8	4	9	1	9

padding →

0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0
0	0	3	5	9	1	10	0	0
0	0	13	2	4	6	11	0	0
0	0	16	24	9	13	1	0	0
0	0	7	1	6	8	3	0	0
0	0	8	4	9	1	9	0	0
0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0

In Keras

- padding = "valid": no padding
- padding = "same": padding with 0 evenly left/right, up/down
- padding = "same" with strides = (1,1): output has same size as input

(W, H)

(W + 2P, H + 2P)

1. Convolution Neural Network

Hyper-Parameters of Conv2D :

- Depth (L): 3 or 1
- Number of Filter: F
- kernel: (K, K): (3,3) or (5,5)
- Stride (S, S): (1,1) or (2,2)
- Padding (P)_w

Formulation to compute the output size of a convolutional layer from an image with size (W, H)?

H

0	0	0	0	0	0	...
0	156	155	156	158	158	...
0	153	154	157	159	159	...
0	149	151	155	158	159	...
0	146	146	149	153	158	...
0	145	143	143	148	158	...
...

1. Convolution Neural Network

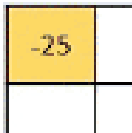
Hyper-Parameters of Conv2D :

- Depth (L): 3 or 1
- Number of Filter: F
- kernel: (K, K): (3,3) or (5,5)
- Stride (S, S): (1,1) or (2,2)
- Padding (P) W_{in}

Formulation to compute the output size of a convolutional layer from an image with size (W, H)?

H_{in}	156	155	156	158	158
	153	154	157	159	159
	149	151	155	158	159
	146	146	149	153	158
	145	143	143	148	158



$$W_{out} = \frac{W_{in} - K + 2P}{S} + 1$$


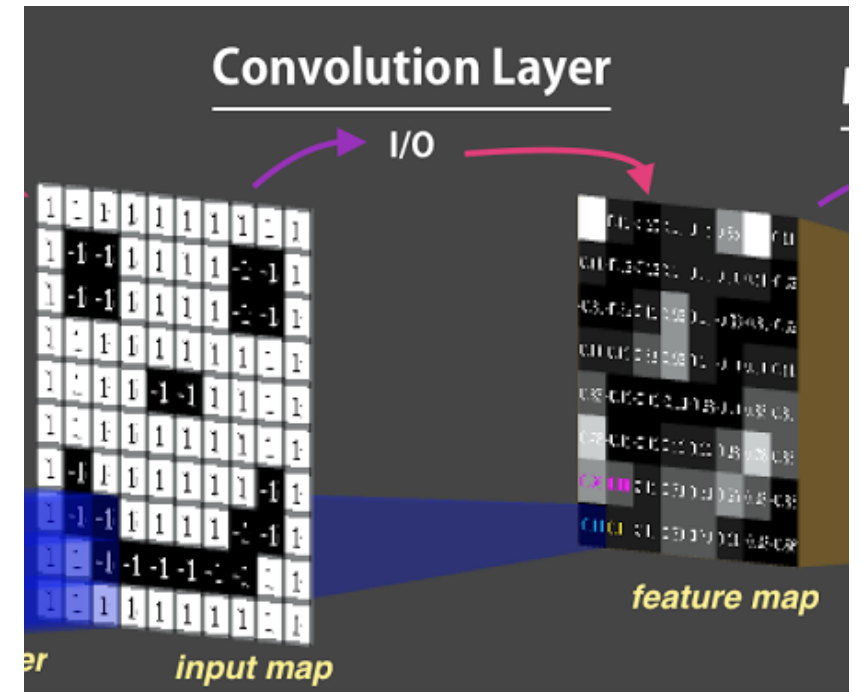
Convolved features

-25				0000
				0000
				0000
				0000
				0000
0000	0000	0000	0000	0000

1. Convolution Neural Network

How to add Conv2D in keras?


```
model = Sequential()  
model.add(Conv2D(F, (K, K), strides=(S, S), activation='relu', padding="same", input_shape=(32, 32, L)))
```



1. Convolution Neural Network

Convolutional Layer (CNN or ConvNet):

- The CNN will reduce the original RGB images to its Convolutional Layer
- Multiple layers can be applied



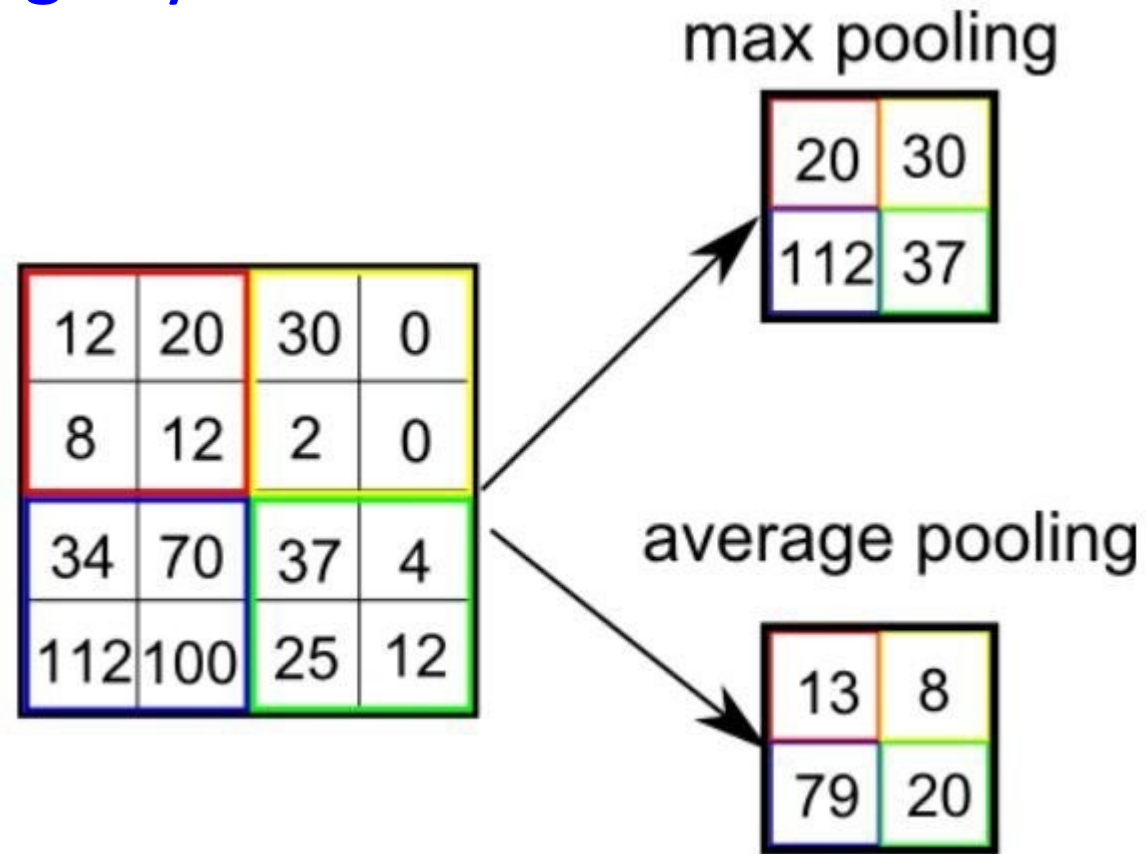
1. Convolution Neural Network

Pooling Layer

- Pooling Layer should follow Convolutional Layer
- Similar to the Convolutional Layer, the Pooling layer is responsible for reducing the spatial size of the Convolved Feature.
- This is to decrease the computational power required to process the data through dimensionality reduction
- Two types of Pooling: Max Pooling & Average Pooling.

1. Convolution Neural Network

Pooling Layer




In which Max Pooling performs a lot better than Average Pooling.

1. Convolution Neural Network

Pooling Layer


- The image after Max Pooling layer would look like:



1. Convolution Neural Network

Flatten Layer

- Once the images have passed through Convolution Layer and Pooling Layer, its size has been reduced greatly and ready for MLP training (or to another Convolution steps).
- The image is then flatten to a column vector and passed through feed-forward NN and BackPropagation applied to every iteration.
- Softmax activation function is applied to classified the multi-output/multi-labels



1. Convolution Neural Network

Some other useful layers?

1. Convolution Neural Network

Batch Normalization

- A process to make Deep neural networks faster and more stable through adding extra layers in a deep neural network.
- The new layer performs the standardizing and normalizing operations on the input of a layer coming from a previous layer.
- Normalize the amounts of weights trained between layers during training
- It usually goes after Conv2D layers or after Dense layer

```
from tensorflow.keras.layers BatchNormalization
model = Sequential()
model.add(Conv2D(75, (3, 3), strides=1, activation="relu", input_shape=(28, 28, 1)))
model.add(BatchNormalization())
model.add(MaxPool2D((2, 2), strides=2))
```

1. Convolution Neural Network


Dropout

- Dropout is a regularization method that approximates training a large number of neural networks with different architectures in parallel.
- Dropout helps to avoid Overfitting
- Dropout is implemented per layer in the NN
- Dropout is not used after training when making a prediction with the fit network.

```
from tensorflow.keras.layers import BatchNormalization
model = Sequential()
model.add(Conv2D(75, (3, 3), strides=1, activation="relu", input_shape=(28, 28, 1)))
model.add(Dropout(0.2))
```

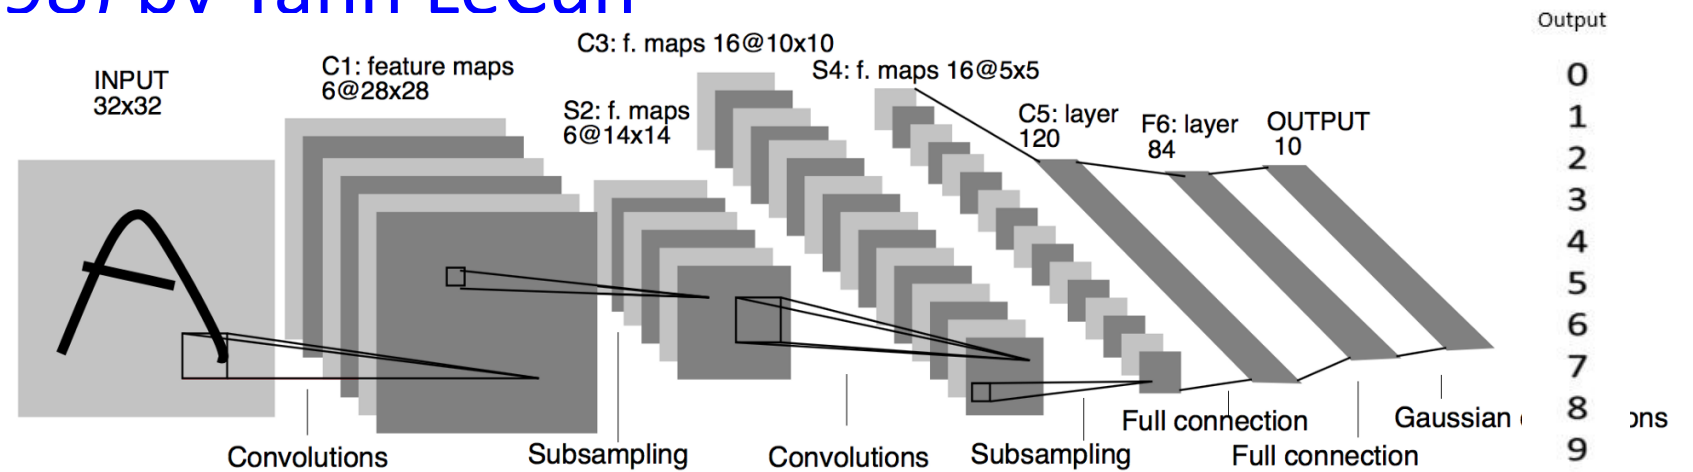
1. Convolution Neural Network

A Sample of CNN:



1. Convolution Neural Network


Lenet-5 (1998) by Yann LeCun



- LeNet-5 is designed for handwritten and machine-printed character recognition
- Input of 32x32x1
- Total parameters: 60k
- Activation function: tanh

1. Convolution Neural Network

Lenet-5 (1998) by Yann LeCun

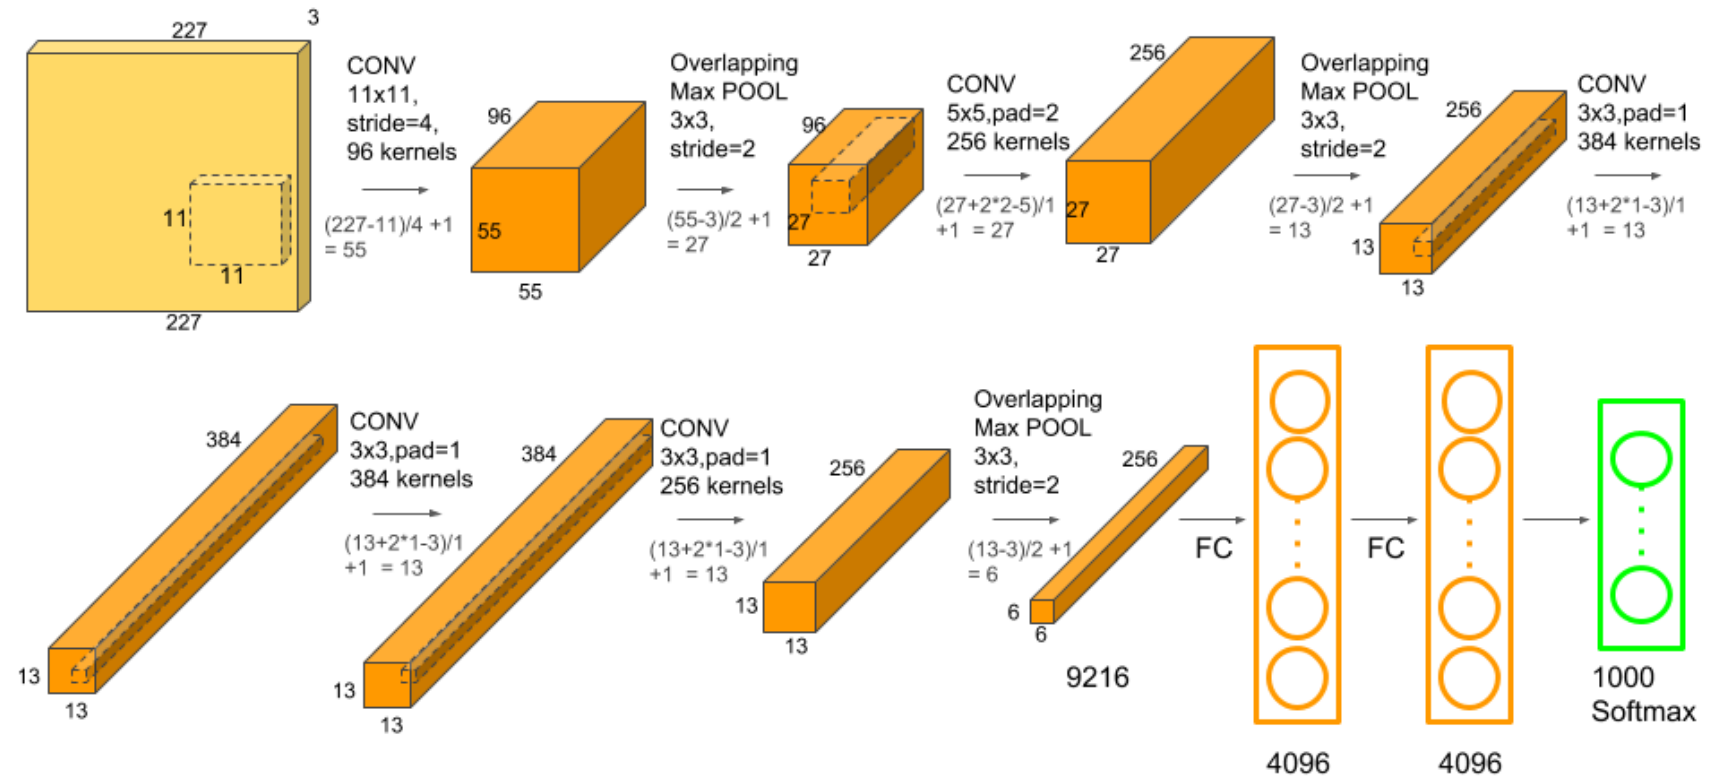


```
model = Sequential()
model.add(Conv2D(6, (5, 5), strides=(1, 1), activation='tanh', padding="valid", input_shape=(32, 32, 1)))
model.add(AveragePooling2D(pool_size=(2, 2), strides=(2, 2)))
model.add(Conv2D(16, (5, 5), strides=(1, 1), activation='tanh', padding="valid"))
model.add(AveragePooling2D(pool_size=(2, 2), strides=(2, 2)))
model.add(Conv2D(120, (5, 5), strides=(1, 1), activation='tanh', padding="valid"))
model.add(Flatten())
model.add(Dense(84, activation='tanh'))
model.add(Dense(10, activation='softmax'))
```


1. Convolution Neural Network

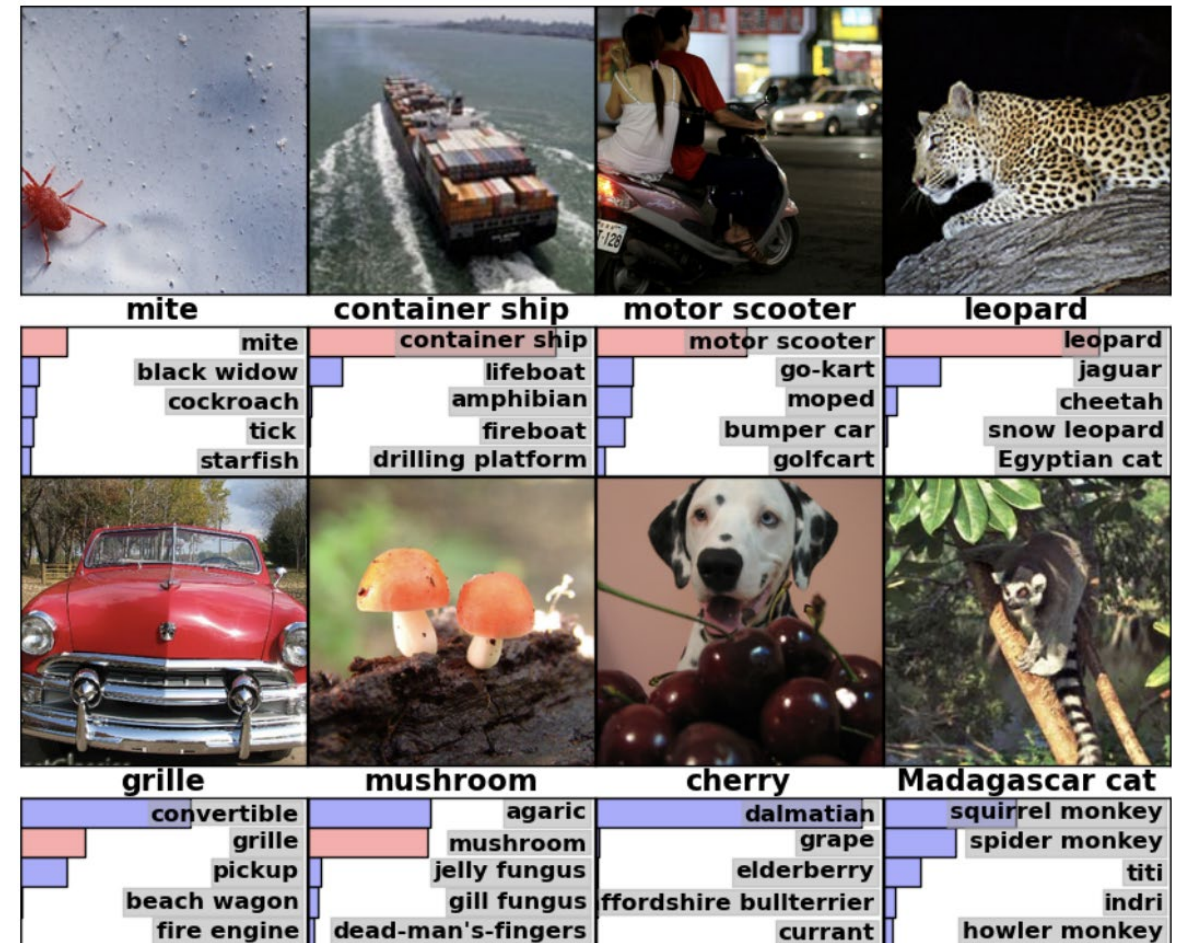
AlexNet (2012) by Hinton Alex Krizhevsky

- AlexNet won the 2012 ImageNet challenge
- Input of 227x227x3
- Total parameters: 60M
- Activation: ReLU



1. Convolution Neural Network


AlexNet (2012) by Hinton Alex Krizhevsky



1. Convolution Neural Network

VGG16 (2014) - Visual Geometry Group


- VGG16 runner up of 2014 ImageNet challenge
- 16 layers: 13 ConvNet, 3 Fully Connected
- Total Parameters: 130M



1. Convolution Neural Network

GoogleNet (2014)

- GoogleNet won the 2014 ImageNet challenge
- Introduced Inception Network
- 22 layers deep with 27 pooling layers and 9 inception models




2. Recurrent Neural Network

2. Recurrent Neural Network

Introduction

- RNNs are type of Deep Learning models with built-in feedback mechanism.
- The output of a particular layer can be **re-fed** as the input in order to predict the output.
- This is different from traditional ML where output/predictand cannot be used as input




Recurrent Neural Networks have loops.

2. Recurrent Neural Network

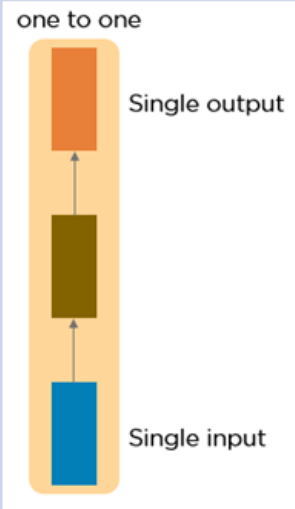
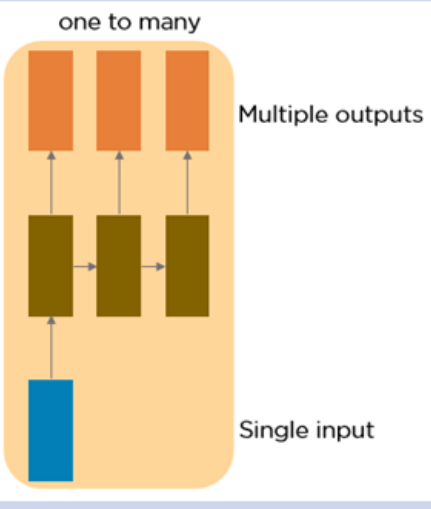
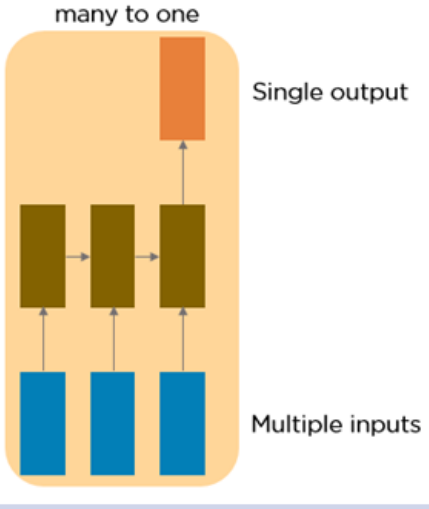
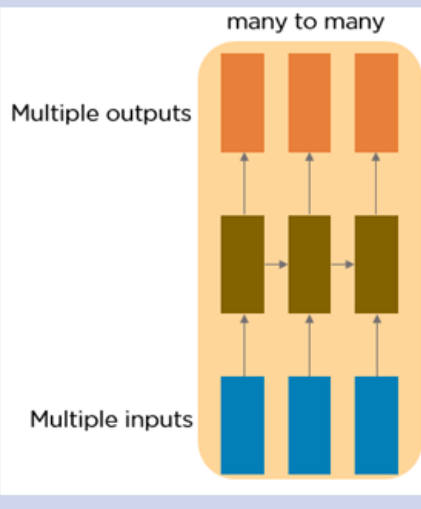


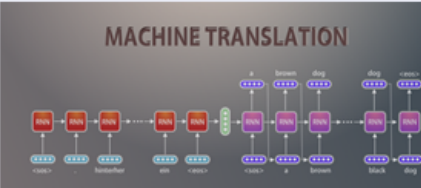
Introduction

- Unroll the RNN loop



2. Recurrent Neural Network

Type of RNNs

One to One	One to Many	Many to One	Many to Many
 <p>one to one</p> <p>Single output</p> <p>Single input</p>	 <p>one to many</p> <p>Multiple outputs</p> <p>Single input</p>	 <p>many to one</p> <p>Single output</p> <p>Multiple inputs</p>	 <p>many to many</p> <p>Multiple outputs</p> <p>Multiple inputs</p>
<p>So called Vanilla NN. Similar to Backpropagation for general ML problem</p>	<p>Image captioning</p>  <p>IMAGE → CAPTION</p> <ul style="list-style-type: none">• large brown dog running away from the sprinkler in the grass• a brown dog chases the water from a sprinkler on a lawn• a brown dog running on a lawn near a garden hose• a brown dog plays with the hose• it is about to interact with a fence	<p>Sentiment analysis</p>  <p>SENTIMENT POSITIVE</p> <p>SENTIMENT NEGATIVE</p>	<p>Machine translation</p>  <p>MACHINE TRANSLATION</p>

2. Recurrent Neural Network

Applications

It is specifically designed for Sequential problem **Weather forecast, Stock forecast, Image captioning, Natural Language Processing, Speech/Voice Recognition**

2. Recurrent Neural Network

Some Disadvantages of RNN:

- Computationally Expensive and large memory requested
- RNN is sensitive to changes in parameters and having gradient problem such as **Exploding Gradient** or **Vanishing Gradient**
- In order to resolve the gradient problem of RNN, a method **Long-Short Term Memory (LSTM)** is proposed.

In this limited workshop, we only cover LSTM for timeseries forecast problem (stock forecast and weather forecast)

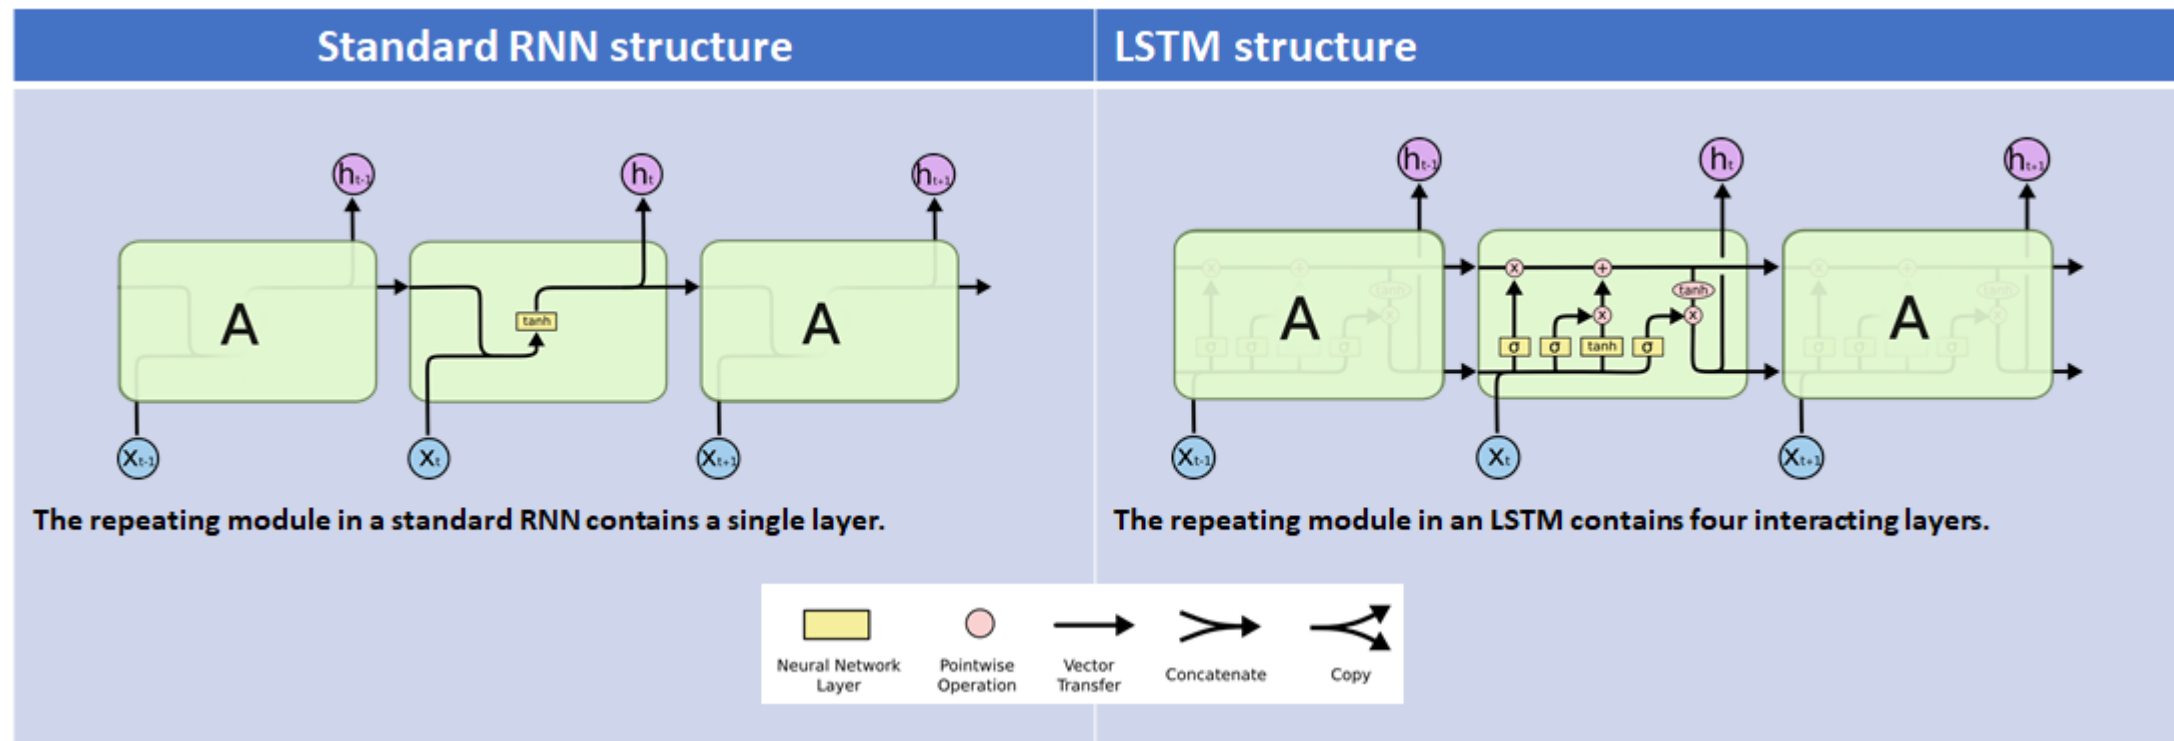
2. Recurrent Neural Network

Long-Short Term Memory model - LSTM

- LSTMs are a special kind of RNN — capable of learning long-term dependencies by remembering information for long periods is the default behavior.
- They were introduced by Hochreiter & Schmidhuber (1997) and were refined and popularized by many people
- LSTMs are explicitly designed to avoid the long-term dependency problem.

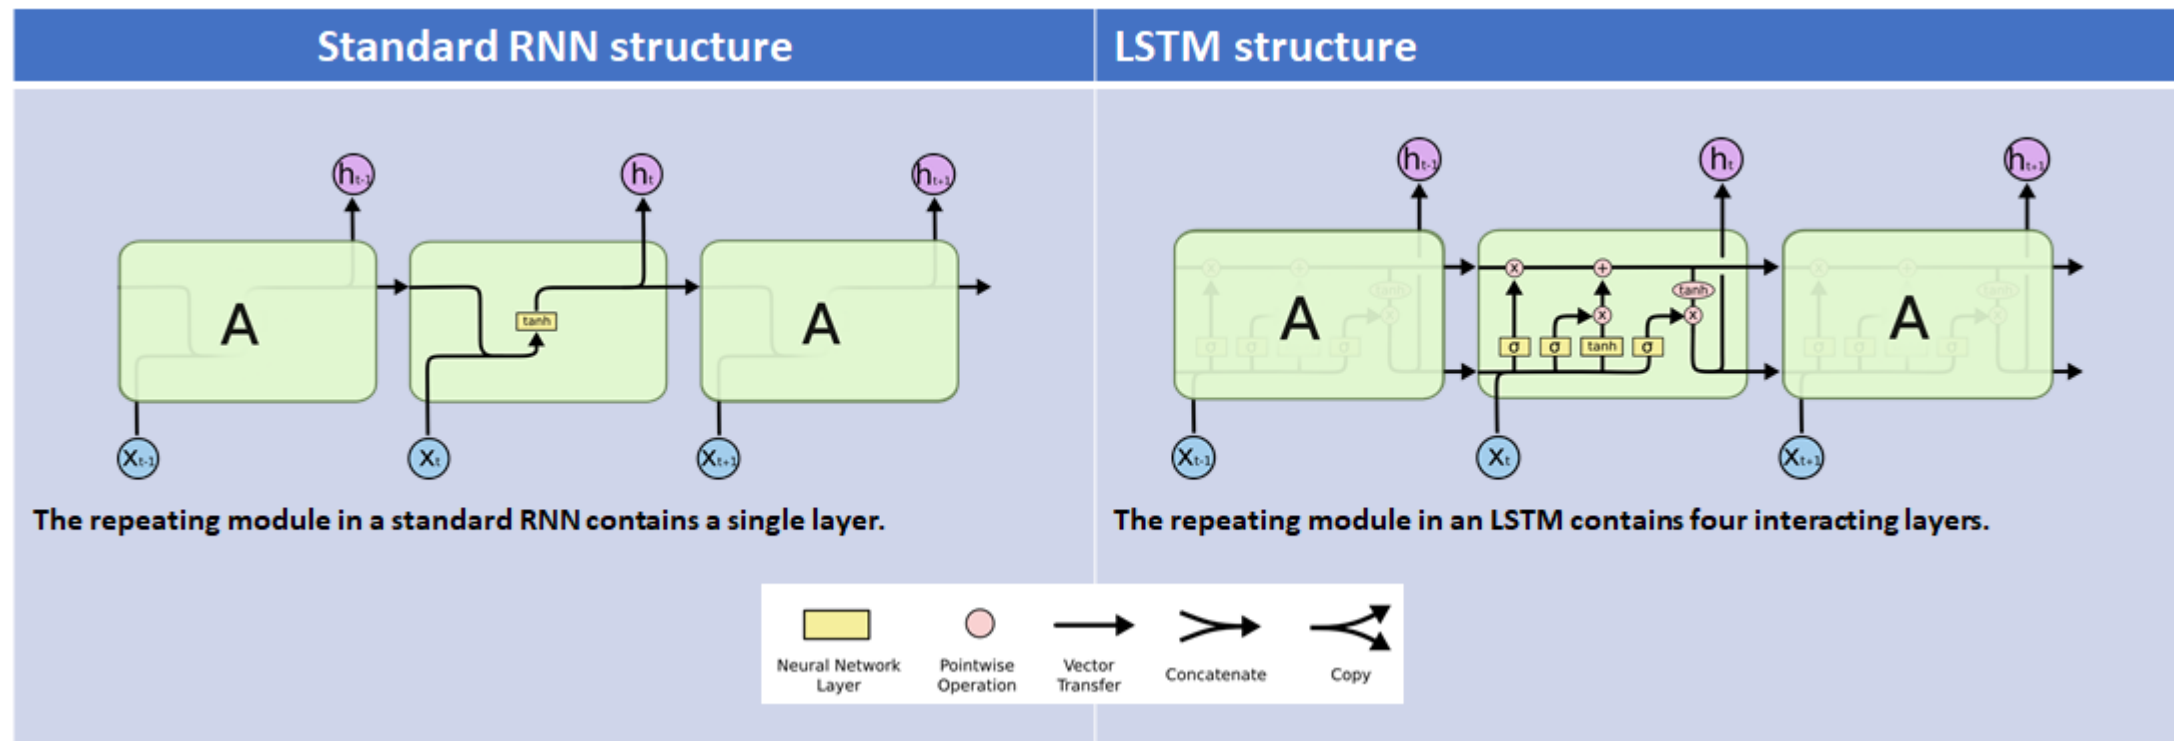
2. Recurrent Neural Network

Comparison between traditional RNN and LSTM



2. Recurrent Neural Network

Comparison between traditional RNN and LSTM



```
model = Sequential()  
model.add(LSTM(128, return_sequences=True, input_shape= (x_train.shape[1], 1)))  
model.add(LSTM(64, return_sequences=False))
```

2. Recurrent Neural Network

Hands-on section